

# Machine Learning Assignment Lab - 2

## Naive Bayes and KNN

Keshav K S 3122 23 5001 067

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import time
import math
from sklearn.model_selection import train_test_split, GridSearchCV,
RandomizedSearchCV, cross_val_score
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.naive_bayes import GaussianNB, MultinomialNB, BernoulliNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import (accuracy_score, precision_score,
recall_score, f1_score, confusion_matrix, roc_curve, auc,
roc_auc_score)

print("Libraries imported successfully.")
Libraries imported successfully.

print("Step 1 : Loading the dataset")

url =
"https://archive.ics.uci.edu/ml/machine-learning-databases/spambase/
spambase.data"
column_names = [
    "word_freq_make", "word_freq_address", "word_freq_all",
"word_freq_3d", "word_freq_our",
    "word_freq_over", "word_freq_remove", "word_freq_internet",
"word_freq_order", "word_freq_mail",
    "word_freq_receive", "word_freq_will", "word_freq_people",
"word_freq_report", "word_freq_addresses",
    "word_freq_free", "word_freq_business", "word_freq_email",
"word_freq_you", "word_freq_credit",
    "word_freq_your", "word_freq_font", "word_freq_000",
"word_freq_money", "word_freq_hp",
    "word_freq_hpl", "word_freq_george", "word_freq_650",
"word_freq_lab", "word_freq_labs",
    "word_freq_telnet", "word_freq_857", "word_freq_data",
"word_freq_415", "word_freq_85",
```

```

    "word_freq_technology", "word_freq_1999", "word_freq_parts",
"word_freq_pm", "word_freq_direct",
    "word_freq_cs", "word_freq_meeting", "word_freq_original",
"word_freq_project", "word_freq_re",
    "word_freq_edu", "word_freq_table", "word_freq_conference",
"char_freq_;", "char_freq_(", "char_freq_[", "char_freq_!", "char_freq_$", "char_freq_#", "capital_run_length_average",
    "capital_run_length_longest", "capital_run_length_total", "target"
]

```

```

df = pd.read_csv(url, header=None, names=column_names)
print(df.head())
print(df.describe())
print(df.info())

```

Step 1 : Loading the dataset

	word_freq_make	word_freq_address	word_freq_all	word_freq_3d	\
0	0.00	0.64	0.64	0.0	
1	0.21	0.28	0.50	0.0	
2	0.06	0.00	0.71	0.0	
3	0.00	0.00	0.00	0.0	
4	0.00	0.00	0.00	0.0	

	word_freq_our	word_freq_over	word_freq_remove	word_freq_internet	\
0	0.32	0.00	0.00	0.00	
1	0.14	0.28	0.21	0.07	
2	1.23	0.19	0.19	0.12	
3	0.63	0.00	0.31	0.63	
4	0.63	0.00	0.31	0.63	

	word_freq_order	word_freq_mail	...	char_freq_;	char_freq_(	\
0	0.00	0.00	...	0.00	0.000	
1	0.00	0.94	...	0.00	0.132	
2	0.64	0.25	...	0.01	0.143	
3	0.31	0.63	...	0.00	0.137	
4	0.31	0.63	...	0.00	0.135	

	char_freq_[	char_freq_!	char_freq_\$	char_freq_#	\
0	0.0	0.778	0.000	0.000	
1	0.0	0.372	0.180	0.048	
2	0.0	0.276	0.184	0.010	
3	0.0	0.137	0.000	0.000	
4	0.0	0.135	0.000	0.000	

```

    capital_run_length_average  capital_run_length_longest \
0                  3.756                   61
1                  5.114                   101
2                  9.821                  485
3                  3.537                   40
4                  3.537                   40

    capital_run_length_total  target
0                  278        1
1                 1028        1
2                 2259        1
3                 191         1
4                 191         1

[5 rows x 58 columns]
      word_freq_make  word_freq_address  word_freq_all  word_freq_3d
\
count      4601.000000          4601.000000          4601.000000          4601.000000
mean       0.104553           0.213015           0.280656           0.065425
std        0.305358           1.290575           0.504143           1.395151
min        0.000000           0.000000           0.000000           0.000000
25%        0.000000           0.000000           0.000000           0.000000
50%        0.000000           0.000000           0.000000           0.000000
75%        0.000000           0.000000           0.420000           0.000000
max        4.540000          14.280000          5.100000          42.810000

      word_freq_our  word_freq_over  word_freq_remove
word_freq_internet \
count      4601.000000          4601.000000          4601.000000
4601.000000
mean       0.312223           0.095901           0.114208
0.105295
std        0.672513           0.273824           0.391441
0.401071
min        0.000000           0.000000           0.000000
0.000000
25%        0.000000           0.000000           0.000000
0.000000
50%        0.000000           0.000000           0.000000
0.000000
75%        0.380000           0.000000           0.000000

```

	0.000000				
max	10.000000	5.880000		7.270000	
11.110000					
	word_freq_order	word_freq_mail	...	char_freq_;	char_freq_(
\count	4601.000000	4601.000000	...	4601.000000	4601.000000
mean	0.090067	0.239413	...	0.038575	0.139030
std	0.278616	0.644755	...	0.243471	0.270355
min	0.000000	0.000000	...	0.000000	0.000000
25%	0.000000	0.000000	...	0.000000	0.000000
50%	0.000000	0.000000	...	0.000000	0.065000
75%	0.000000	0.160000	...	0.000000	0.188000
max	5.260000	18.180000	...	4.385000	9.752000
	char_freq_[	char_freq_!	char_freq_\$	char_freq_#	\
count	4601.000000	4601.000000	4601.000000	4601.000000	
mean	0.016976	0.269071	0.075811	0.044238	
std	0.109394	0.815672	0.245882	0.429342	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.315000	0.052000	0.000000	
max	4.081000	32.478000	6.003000	19.829000	
	capital_run_length_average	capital_run_length_longest	\		
count	4601.000000		4601.000000		
mean	5.191515		52.172789		
std	31.729449		194.891310		
min	1.000000		1.000000		
25%	1.588000		6.000000		
50%	2.276000		15.000000		
75%	3.706000		43.000000		
max	1102.500000		9989.000000		
	capital_run_length_total	target			
count	4601.000000	4601.000000			
mean	283.289285	0.394045			
std	606.347851	0.488698			
min	1.000000	0.000000			
25%	35.000000	0.000000			
50%	95.000000	0.000000			

75%	266.000000	1.000000
max	15841.000000	1.000000

[8 rows x 58 columns]

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 4601 entries, 0 to 4600

Data columns (total 58 columns):

#	Column	Non-Null Count	Dtype
0	word_freq_make	4601	float64
1	word_freq_address	4601	float64
2	word_freq_all	4601	float64
3	word_freq_3d	4601	float64
4	word_freq_our	4601	float64
5	word_freq_over	4601	float64
6	word_freq_remove	4601	float64
7	word_freq_internet	4601	float64
8	word_freq_order	4601	float64
9	word_freq_mail	4601	float64
10	word_freq_receive	4601	float64
11	word_freq_will	4601	float64
12	word_freq_people	4601	float64
13	word_freq_report	4601	float64
14	word_freq_addresses	4601	float64
15	word_freq_free	4601	float64
16	word_freq_business	4601	float64
17	word_freq_email	4601	float64
18	word_freq_you	4601	float64
19	word_freq_credit	4601	float64
20	word_freq_your	4601	float64
21	word_freq_font	4601	float64
22	word_freq_000	4601	float64
23	word_freq_money	4601	float64
24	word_freq_hp	4601	float64
25	word_freq_hpl	4601	float64
26	word_freq_george	4601	float64
27	word_freq_650	4601	float64
28	word_freq_lab	4601	float64
29	word_freq_labs	4601	float64
30	word_freq_telnet	4601	float64
31	word_freq_857	4601	float64
32	word_freq_data	4601	float64
33	word_freq_415	4601	float64
34	word_freq_85	4601	float64
35	word_freq_technology	4601	float64
36	word_freq_1999	4601	float64
37	word_freq_parts	4601	float64
38	word_freq_pm	4601	float64
39	word_freq_direct	4601	float64

```
40 word_freq_cs           4601 non-null   float64
41 word_freq_meeting      4601 non-null   float64
42 word_freq_original      4601 non-null   float64
43 word_freq_project      4601 non-null   float64
44 word_freq_re            4601 non-null   float64
45 word_freq_edu           4601 non-null   float64
46 word_freq_table          4601 non-null   float64
47 word_freq_conference    4601 non-null   float64
48 char_freq_              4601 non-null   float64
49 char_freq_(             4601 non-null   float64
50 char_freq_[             4601 non-null   float64
51 char_freq_!             4601 non-null   float64
52 char_freq_$             4601 non-null   float64
53 char_freq_#             4601 non-null   float64
54 capital_run_length_average 4601 non-null   float64
55 capital_run_length_longest 4601 non-null   int64
56 capital_run_length_total   4601 non-null   int64
57 target                  4601 non-null   int64
dtypes: float64(55), int64(3)
memory usage: 2.0 MB
None
```

```
print(df.shape)
```

```
(4601, 58)
```

```
print("Step 2 : Performing Data Preprocessing")
```

```
print("Checking Missing Values")
```

```
missing_count = df.isnull().sum()
print("Missing values are : ",missing_count)
```

```
Step 2 : Performing Data Preprocessing
```

```
Checking Missing Values
```

```
Missing values are : word_freq_make          0
word_freq_address          0
word_freq_all               0
word_freq_3d                0
word_freq_our                0
word_freq_over                0
word_freq_remove                0
word_freq_internet                0
word_freq_order                0
word_freq_mail                0
word_freq_receive                0
word_freq_will                0
word_freq_people                0
word_freq_report                0
```

```
word_freq_addresses      0
word_freq_free           0
word_freq_business        0
word_freq_email           0
word_freq_you              0
word_freq_credit           0
word_freq_your             0
word_freq_font             0
word_freq_000              0
word_freq_money             0
word_freq_hp                0
word_freq_hpl               0
word_freq_george            0
word_freq_650               0
word_freq_lab                0
word_freq_labs               0
word_freq_telnet              0
word_freq_857               0
word_freq_data               0
word_freq_415               0
word_freq_85                 0
word_freq_technology          0
word_freq_1999               0
word_freq_parts               0
word_freq_pm                  0
word_freq_direct              0
word_freq_cs                  0
word_freq_meeting              0
word_freq_original             0
word_freq_project              0
word_freq_re                  0
word_freq_edu                  0
word_freq_table                 0
word_freq_conference             0
char_freq_';                   0
char_freq_('                     0
char_freq_[                     0
char_freq_!                     0
char_freq_$                     0
char_freq_#                     0
capital_run_length_average       0
capital_run_length_longest        0
capital_run_length_total          0
target                           0
dtype: int64
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
import pandas as pd
import numpy as np

corr_matrix = df.corr()

target_corr = abs(corr_matrix["target"])
top_features = target_corr[target_corr > 0.25].index
df_top = df[top_features]

plt.rcParams.update({
    "font.size": 15,
    "font.family": "Arial"
})

bxwidth = 1

fig, ax = plt.subplots(figsize=(12, 10))

mask = np.triu(np.ones_like(df_top.corr()), dtype=bool)

sns.heatmap(
    df_top.corr(),
    annot=True,
    fmt=".2f",
    cmap="coolwarm",
    cbar=True,
    mask=mask,
    ax=ax,
    linewidths=0.5,
    linecolor='black'
)

ax.set_title("")

for spine in ax.spines.values():
    spine.set_linewidth(bxwidth)

fig.text(
    0.5, 0.01,
    "Heatmap showing features with high correlation to the target variable.",
    ha="center",
    fontsize=15
)

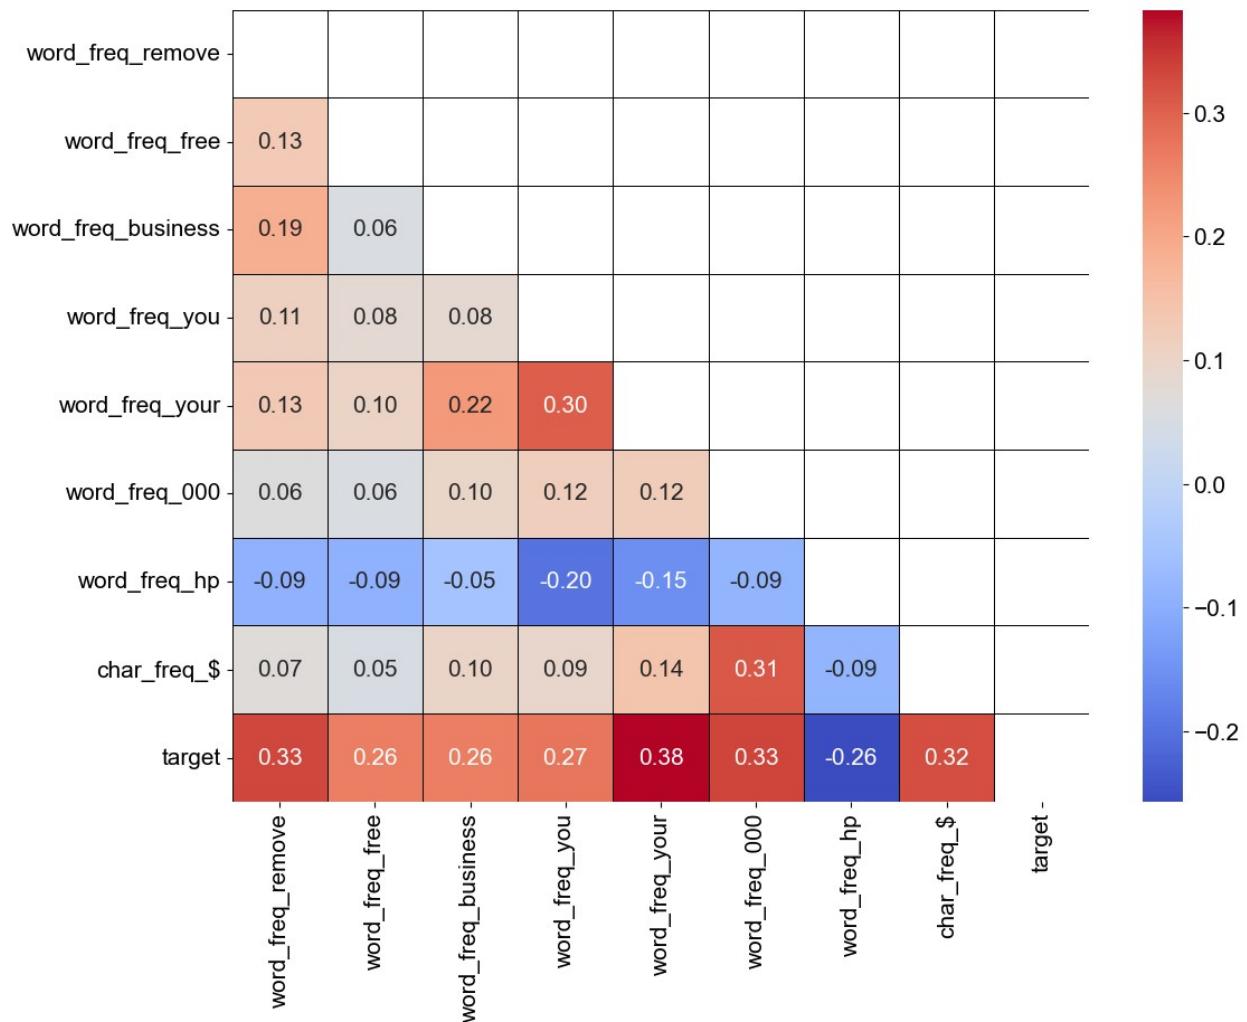
plt.tight_layout(rect=[0, 0.05, 1, 1])
```

```

plt.savefig(
    "correlation_heatmap.eps",
    format="eps",
    dpi=600,
    bbox_inches="tight"
)

plt.show()

```



Heatmap showing features with high correlation to the target variable.

```

print("Step 4 : Visualizing class distribution and feature behavior")

numeric_df =
df.drop(columns=['target']).select_dtypes(include=[ "int64",
"float64"])
cols = numeric_df.columns

```

```

plt.rcParams.update({
    "font.size": 15,
    "font.family": "Arial"
})

bxwidth = 1
rows = math.ceil(len(cols) / 3)

fig, axes = plt.subplots(rows, 3, figsize=(15, 5 * rows))
axes = axes.flatten()

subplot_labels = []
for i in range(len(cols)):
    if i < 26:
        subplot_labels.append(f"chr(97+i))")
    else:
        subplot_labels.append(f"chr(97+(i//26)-1){chr(97+(i%26))}"))

for i, col in enumerate(cols):
    ax = axes[i]

    data_non_spam = df[df['target'] == 0][col].dropna()
    data_spam = df[df['target'] == 1][col].dropna()

    ax.hist(
        [data_non_spam, data_spam],
        bins=20,
        edgecolor="black",
        label=['Non-Spam', 'Spam'],
        color=['#1f77b4', '#ff7f0e'],
        stacked=True
    )

    ax.set_title("")

    ax.set_xlabel(col)
    ax.set_ylabel("Frequency")

    ax.legend(prop={'size': 10})

    ax.text(
        0.5, -0.30,
        f"{subplot_labels[i]} {col} vs target",
        ha="center",
        va="top",
        transform=ax.transAxes,
        fontsize=13
    )

```

```
for spine in ax.spines.values():
    spine.set_linewidth(bxwidth)

for j in range(i + 1, len(axes)):
    fig.delaxes(axes[j])

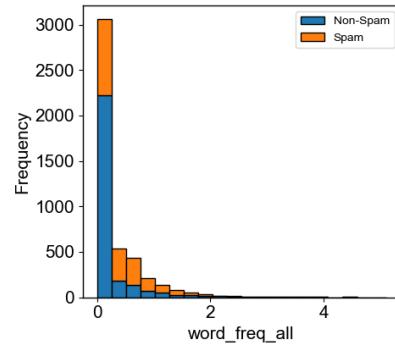
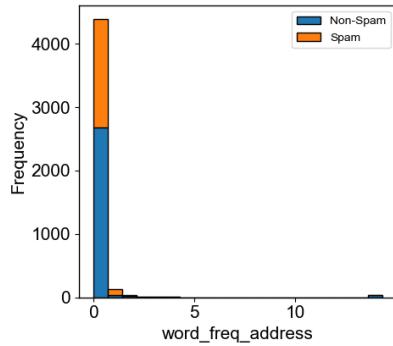
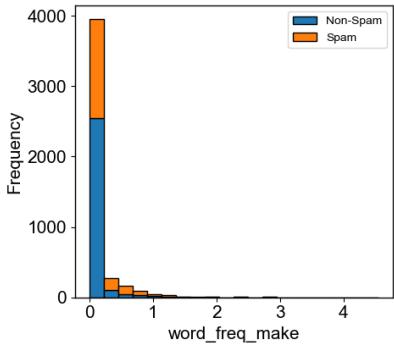
fig.text(
    0.5, 0.005,
    "Histograms showing the distribution of numeric features separated
    by target class.",
    ha="center",
    fontsize=15
)

plt.tight_layout(rect=[0, 0.02, 1, 1])

plt.savefig(
    "histogram_numeric_features.png",
    dpi=600,
    bbox_inches="tight"
)

plt.show()

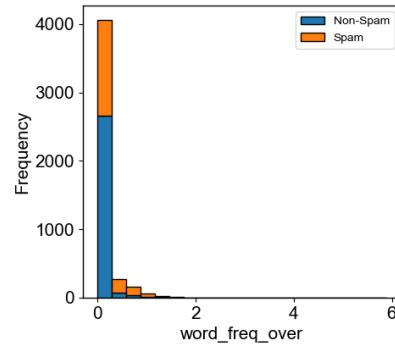
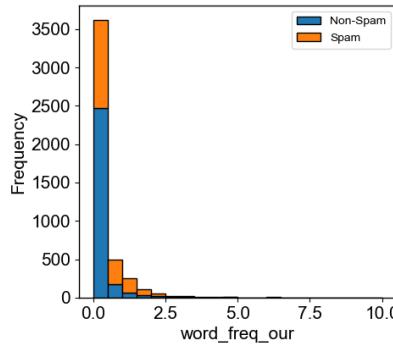
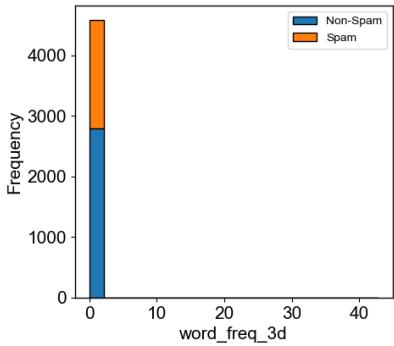
Step 4 : Visualizing class distribution and feature behavior
```



a) word\_freq\_make vs target

b) word\_freq\_address vs target

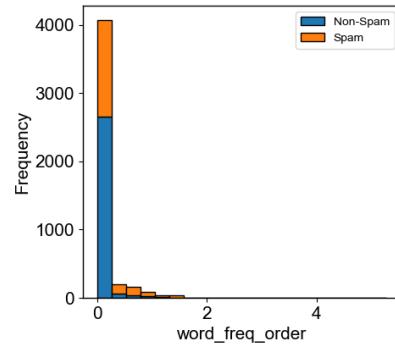
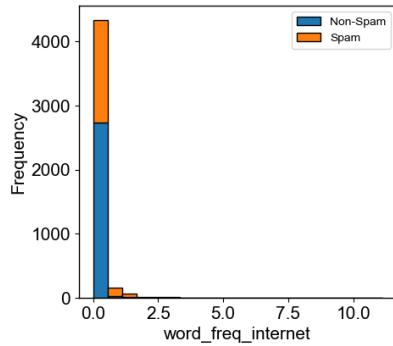
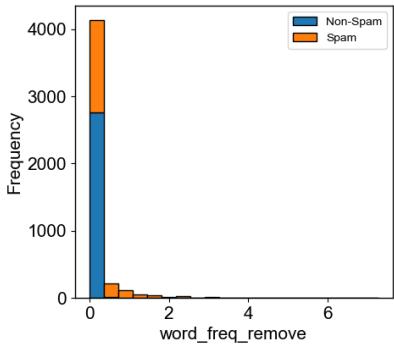
c) word\_freq\_all vs target



d) word\_freq\_3d vs target

e) word\_freq\_our vs target

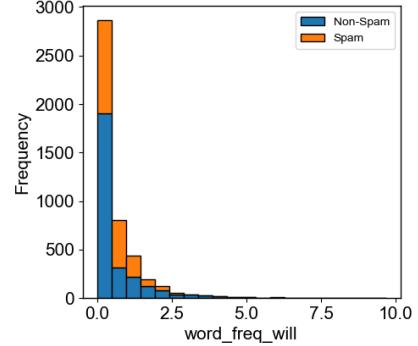
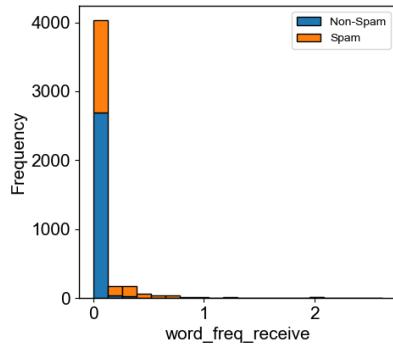
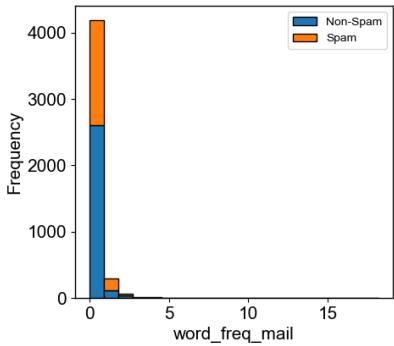
f) word\_freq\_over vs target



g) word\_freq\_remove vs target

h) word\_freq\_internet vs target

i) word\_freq\_order vs target



j) word\_freq\_mail vs target

k) word\_freq\_receive vs target

l) word\_freq\_will vs target



```

class_counts = df['target'].value_counts().sort_index()
classes = ['Non-Spam', 'Spam']
counts = [class_counts[0], class_counts[1]]

plt.rcParams.update({
    "font.size": 15,
    "font.family": "Arial"
})

bxwidth = 1

fig, ax = plt.subplots(figsize=(6, 6))

bars = ax.bar(classes, counts, color=['#1f77b4', '#ff7f0e'],
edgecolor="black")

ax.set_title("")
ax.set_xlabel("Class")
ax.set_ylabel("Count")

for bar in bars:
    height = bar.get_height()
    ax.text(bar.get_x() + bar.get_width()/2., height,
            f'{int(height)}',
            ha='center', va='bottom', fontsize=13)

for spine in ax.spines.values():
    spine.set_linewidth(bxwidth)

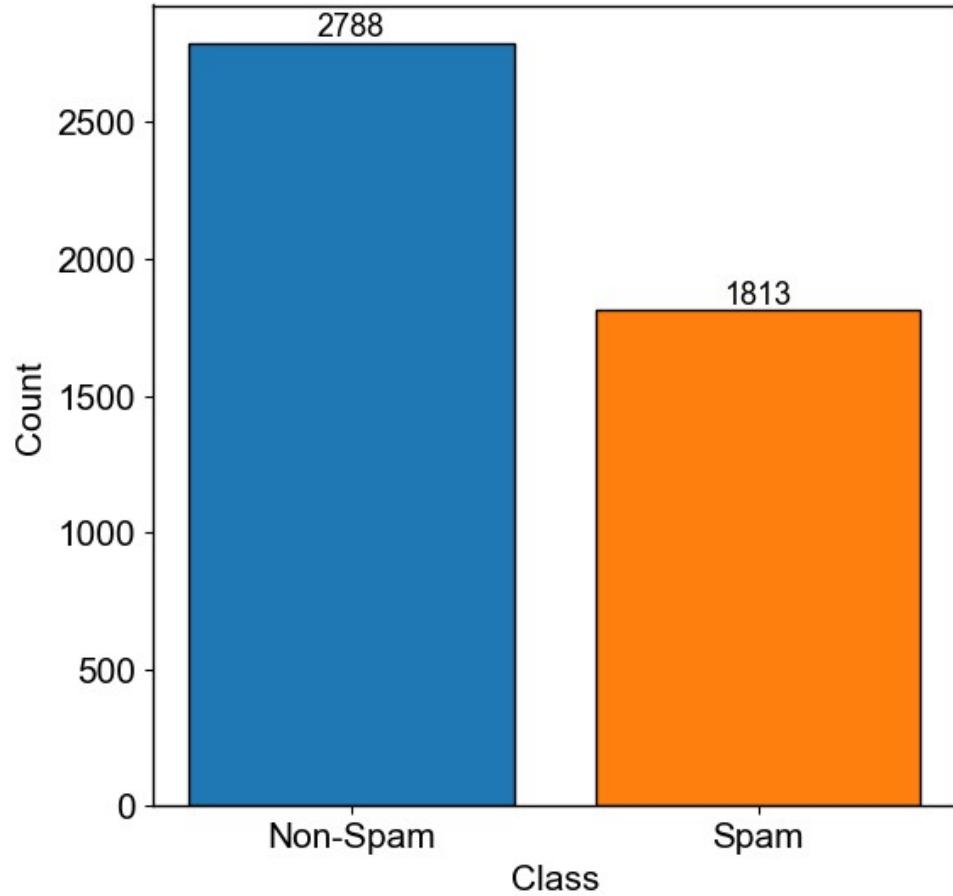
fig.text(
    0.5, -0.05,
    "Class distribution showing the balance between Non-Spam and Spam
emails.",
    ha="center",
    fontsize=15
)

plt.tight_layout(rect=[0, 0.05, 1, 1])

plt.savefig(
    "class_distribution.png",
    dpi=600,
    bbox_inches="tight"
)

plt.show()

```



Class distribution showing the balance between Non-Spam and Spam emails.

```
import matplotlib.pyplot as plt

features = [
    'word_freq_make',
    'word_freq_address',
    'word_freq_all',
    'word_freq_3d',
    'word_freq_our'
]

data = [df[feature] for feature in features]

plt.figure(figsize=(12, 6))

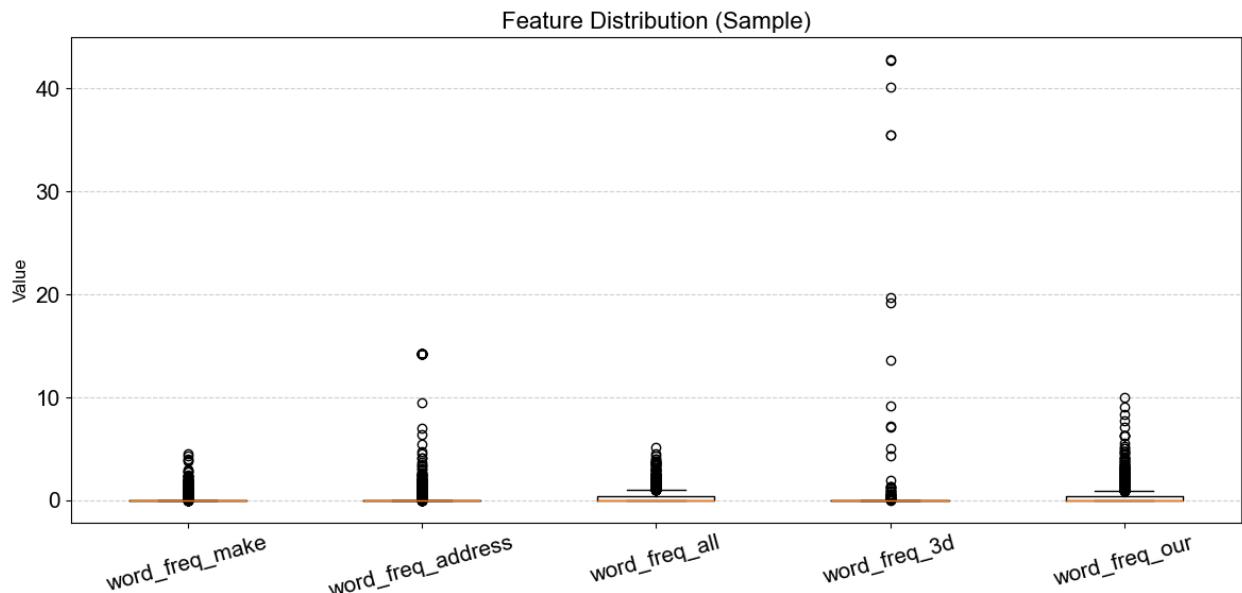
plt.boxplot(
    data,
    labels=features,
    showfliers=True,
```

```

)
plt.title("Feature Distribution (Sample)", fontsize=16)
plt.ylabel("Value", fontsize=12)
plt.xticks(rotation=15)
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.savefig("feature_distribution", dpi = 600, bbox_inches = "tight")
plt.tight_layout()
plt.show()

C:\Users\KESHA\AppData\Local\Temp\ipykernel_2000\2819539952.py:15:
MatplotlibDeprecationWarning: The 'labels' parameter of boxplot() has
been renamed 'tick_labels' since Matplotlib 3.9; support for the old
name will be dropped in 3.11.
    plt.boxplot(

```



```

print("Step 5 : Splitting the dataset into training and testing sets")

X = df.drop(columns=['target'])
y = df['target']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

print("Split Successful.\n")
print(f"Total Samples: {len(df)}")
print(f"Training Set: {X_train.shape[0]} samples (80%)")
print(f"Testing Set: {X_test.shape[0]} samples (20%)")

```

```

print("\nClass distribution in Training Set:")
print(y_train.value_counts(normalize=True))

Step 5 : Splitting the dataset into training and testing sets
Split Successful.

Total Samples: 4601
Training Set: 3680 samples (80%)
Testing Set: 921 samples (20%)

Class distribution in Training Set:
target
0    0.605978
1    0.394022
Name: proportion, dtype: float64

print("Scaling data")

scaler=StandardScaler()
X_train_scaled=scaler.fit_transform(X_train)
X_test_scaled=scaler.transform(X_test)

Scaling data

from sklearn.metrics import classification_report

print("Step 6: Training Naïve Bayes Variants")

start = time.time()
gnb = GaussianNB()
gnb.fit(X_train_scaled, y_train)
time_gnb= time.time() - start
y_pred_gnb = gnb.predict(X_test)

print("The training time for Gaussian Naive Bayes is: ",time_gnb)
print("Accuracy score",accuracy_score(y_test,y_pred_gnb))

print("Precision score",precision_score(y_test,y_pred_gnb))
print("Recall score",recall_score(y_test,y_pred_gnb))
print("F1 score",f1_score(y_test,y_pred_gnb))
print("Confusion matrix",confusion_matrix(y_test,y_pred_gnb))
print("Report",classification_report(y_test,y_pred_gnb))

bnb = BernoulliNB(binarize=0.0)
bnb.fit(X_train, y_train)
y_pred_bnb = bnb.predict(X_test)

```

```

acc_bnb = accuracy_score(y_test, y_pred_bnb)

start = time.time()
bnb = BernoulliNB()
bnb.fit(X_train, y_train)
time_bnb=time.time()-start
y_pred_bnb = bnb.predict(X_test)
print("The training time for Bernoulli Naive Bayes is: ",time_bnb)

print("Accuracy score",accuracy_score(y_test,y_pred_bnb))

print("Precision score",precision_score(y_test,y_pred_bnb))
print("Recall score",recall_score(y_test,y_pred_bnb))
print("F1 score",f1_score(y_test,y_pred_bnb))
print("Confusion matrix",confusion_matrix(y_test,y_pred_bnb))
print("Report",classification_report(y_test,y_pred_bnb))

start = time.time()
mnb = MultinomialNB()
mnb.fit(X_train, y_train)
time_mnb=time.time()-start
y_pred_mnb = mnb.predict(X_test)
print("The training time for Multinomial Naive Bayes is: ",time_mnb)

print("Accuracy score",accuracy_score(y_test,y_pred_mnb))

print("Precision score",precision_score(y_test,y_pred_mnb))
print("Recall score",recall_score(y_test,y_pred_mnb))
print("F1 score",f1_score(y_test,y_pred_mnb))
print("Confusion matrix",confusion_matrix(y_test,y_pred_mnb))
print("Report",classification_report(y_test,y_pred_mnb))

```

#### Step 6: Training Naïve Bayes Variants

The training time for Gaussian Naive Bayes is: 0.006426334381103516

Accuracy score 0.5722041259500543

Precision score 0.4791946308724832

Recall score 0.9834710743801653

F1 score 0.644404332129964

Confusion matrix [[170 388]

[ 6 357]]

Report	precision	recall	f1-score	support
0	0.97	0.30	0.46	558
1	0.48	0.98	0.64	363

	precision	recall	f1-score	support
0	0.97	0.30	0.46	558
1	0.48	0.98	0.64	363

accuracy			0.57	921
macro avg	0.72	0.64	0.55	921
weighted avg	0.77	0.57	0.53	921

The training time for Bernoulli Naive Bayes is: 0.00619053840637207  
Accuracy score 0.8762214983713354  
Precision score 0.8716417910447761  
Recall score 0.8044077134986226  
F1 score 0.836676217765043  
Confusion matrix [[515 43]  
 [ 71 292]]

Report	precision	recall	f1-score	support
0	0.88	0.92	0.90	558
1	0.87	0.80	0.84	363

accuracy			0.88	921
macro avg	0.88	0.86	0.87	921
weighted avg	0.88	0.88	0.88	921

The training time for Multinomial Naive Bayes is:  
0.003495931625366211  
Accuracy score 0.7763300760043431  
Precision score 0.7198879551820728  
Recall score 0.7079889807162535  
F1 score 0.7138888888888889  
Confusion matrix [[458 100]  
 [106 257]]

Report	precision	recall	f1-score	support
0	0.81	0.82	0.82	558
1	0.72	0.71	0.71	363

accuracy			0.78	921
macro avg	0.77	0.76	0.77	921
weighted avg	0.78	0.78	0.78	921

e:\anaconda\Lib\site-packages\sklearn\utils\validation.py:2732:  
UserWarning: X has feature names, but GaussianNB was fitted without  
feature names  
warnings.warn(

```
import matplotlib.pyplot as plt
import numpy as np

def plot_confusion_matrix(cm, title):
    plt.figure()
    plt.imshow(cm)
    plt.title(title)
```

```

plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.xticks(np.arange(cm.shape[0]))
plt.yticks(np.arange(cm.shape[0]))

for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        plt.text(j, i, cm[i, j], ha="center", va="center")

plt.tight_layout()
plt.show()

cm_gnb = confusion_matrix(y_test, y_pred_gnb)

plt.figure(figsize=(5, 4))
sns.heatmap(cm_gnb, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix – Gaussian Naive Bayes")
plt.xlabel("Predicted")
plt.ylabel("Actual")

plt.savefig("confusion_matrix_gaussian_nb.png", dpi=600,
bbox_inches="tight")
plt.show()
plt.close()

cm_bnb = confusion_matrix(y_test, y_pred_bnb)

plt.figure(figsize=(5, 4))
sns.heatmap(cm_bnb, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix – Bernoulli Naive Bayes")
plt.xlabel("Predicted")
plt.ylabel("Actual")

plt.savefig("confusion_matrix_bernoulli_nb.png", dpi=600,
bbox_inches="tight")
plt.show()
plt.close()

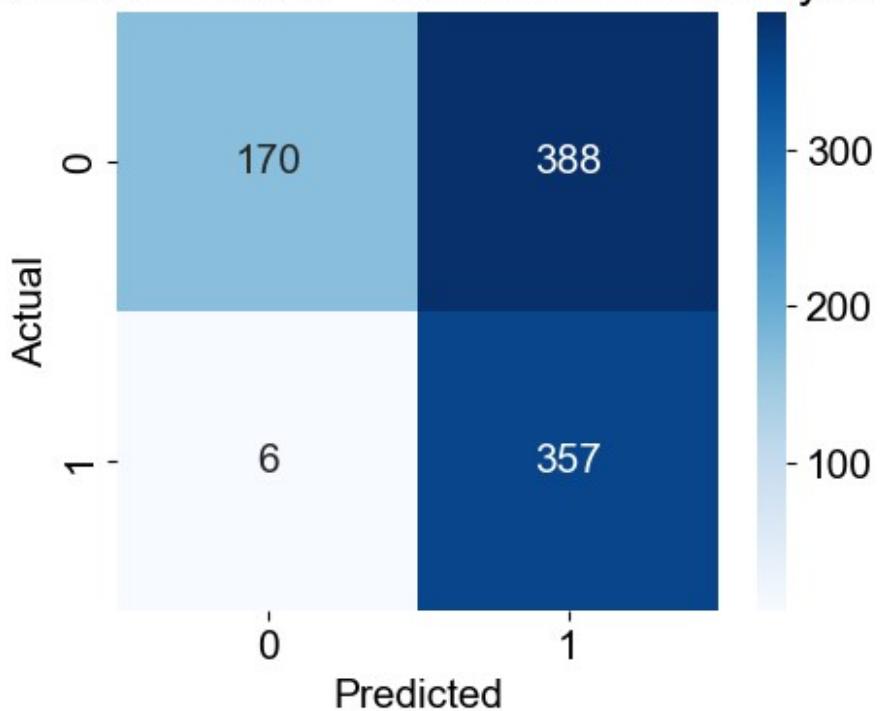
cm_mnb = confusion_matrix(y_test, y_pred_mnb)

plt.figure(figsize=(5, 4))
sns.heatmap(cm_mnb, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix – Multinomial Naive Bayes")
plt.xlabel("Predicted")
plt.ylabel("Actual")

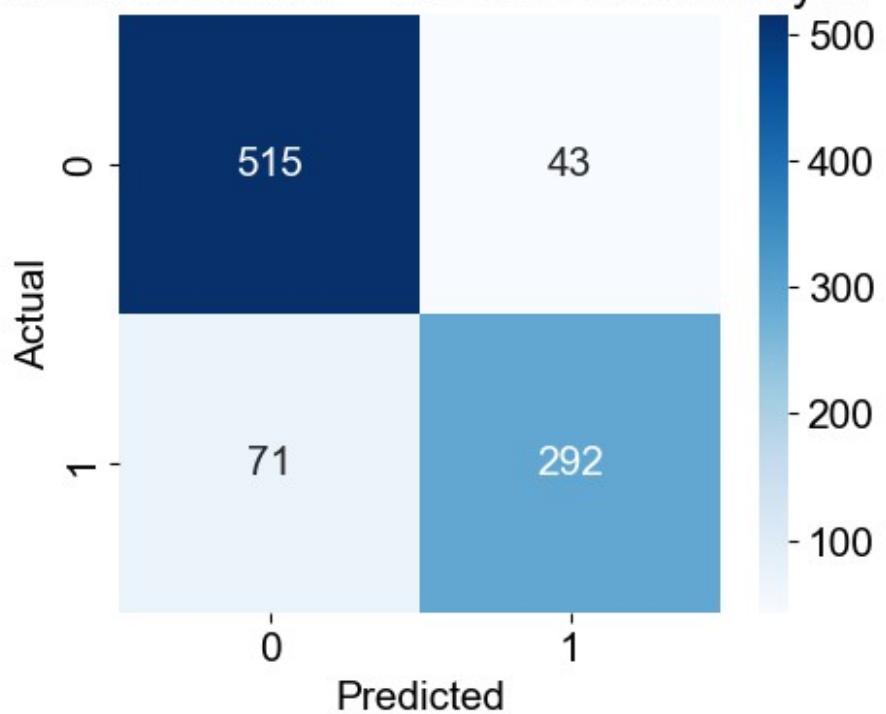
plt.savefig("confusion_matrix_multinomial_nb.png", dpi=600,
bbox_inches="tight")
plt.show()
plt.close()

```

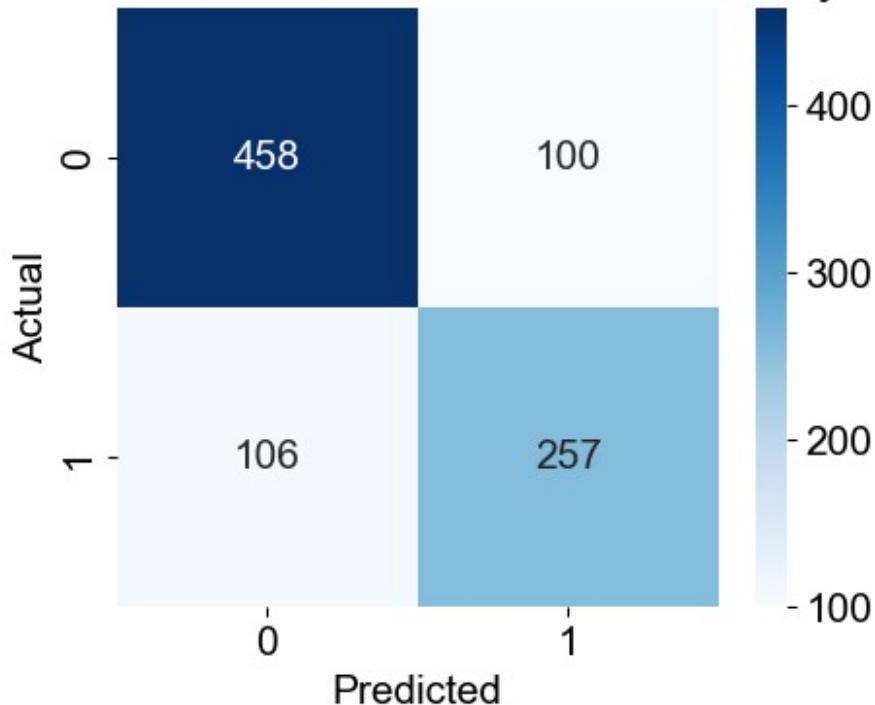
### Confusion Matrix – Gaussian Naive Bayes



### Confusion Matrix – Bernoulli Naive Bayes



## Confusion Matrix – Multinomial Naive Bayes



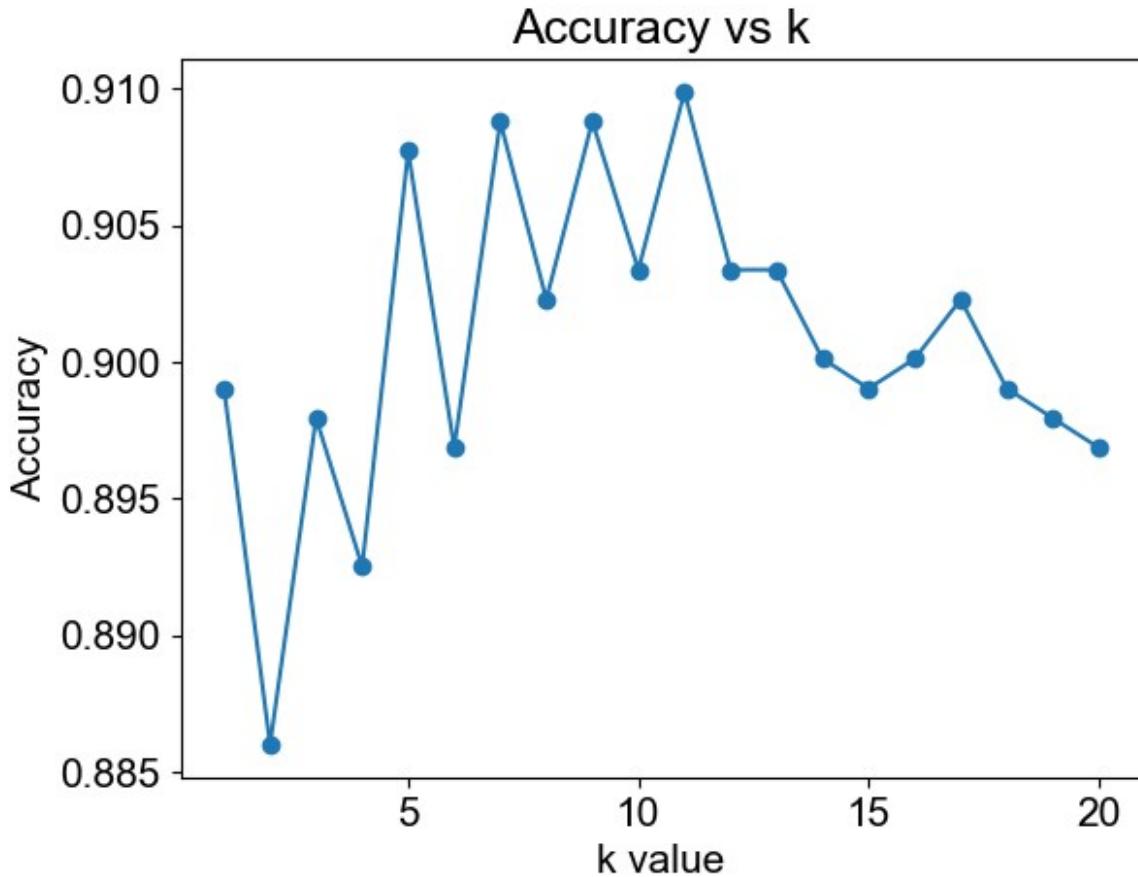
```
print("Step 7: Training Baseline KNN Classifier")

knn = KNeighborsClassifier()
knn.fit(X_train_scaled, y_train)

y_pred_knn = knn.predict(X_test_scaled)

k_values = range(1, 21)
accuracies = []
for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train_scaled, y_train)
    accuracies.append(knn.score(X_test_scaled, y_test))
plt.plot(k_values, accuracies, marker='o')
plt.xlabel('k value')
plt.ylabel('Accuracy')
plt.title('Accuracy vs k')
plt.savefig("accuracy_vs_k", dpi = 600, bbox_inches = "tight")
plt.show()

Step 7: Training Baseline KNN Classifier
```



```
from sklearn.model_selection import StratifiedKFold

skf=StratifiedKFold(n_splits=5,shuffle=True,random_state=42)
cv_scores_base=cross_val_score(
    knn,
    X_train_scaled,
    y_train,
    cv=skf,
    scoring='accuracy'
)
print("Base KNN CV Accuracy:",cv_scores_base.mean())

Base KNN CV Accuracy: 0.8932065217391303

print("Step 8: Hyperparameter Tuning for KNN")

param_grid={
    'n_neighbors':list(range(1,21,2)),
    'weights':[['uniform','distance']],
    'metric':['euclidean','manhattan']
}
```

```

grid=GridSearchCV(
    knn,
    param_grid,
    cv=skf,
    scoring='accuracy',
    n_jobs=-1
)
grid.fit(X_train_scaled,y_train)
print("Grid Best Params:",grid.best_params_)
print("Grid Best CV Accuracy:",grid.best_score_)

Step 8: Hyperparameter Tuning for KNN
Grid Best Params: {'metric': 'manhattan', 'n_neighbors': 9, 'weights': 'distance'}
Grid Best CV Accuracy: 0.9252717391304348

random = RandomizedSearchCV(
    KNeighborsClassifier(),
    param_distributions=param_grid,
    n_iter=5,
    cv=5,
    scoring="accuracy",
    random_state=42
)
random.fit(X_train_scaled, y_train)
print("Random Best Params : ",random.best_params_)
print("Random Best CV Accuracy : ",random.best_score_)

Random Best Params : {'weights': 'distance', 'n_neighbors': 15,
'metric': 'euclidean'}
Random Best CV Accuracy :  0.9209239130434783

best_params=grid.best_params_
knn_final=KNeighborsClassifier(
    n_neighbors=best_params['n_neighbors'],
    weights=best_params['weights'],
    metric=best_params['metric']
)
knn_final.fit(X_train_scaled,y_train)
y_pred_final=knn_final.predict(X_test_scaled)

from sklearn.metrics import confusion_matrix,roc_curve,auc
def compute_metrics(y_true,y_pred):
    cm=confusion_matrix(y_true,y_pred)
    tn,fp,fn,tp=cm.ravel()
    accuracy=(tp+tn)/(tp+tn+fp+fn)
    precision=tp/(tp+fp)
    recall=tp/(tp+fn)

```

```

f1=2*precision*recall/(precision+recall)
specificity=tn/(tn+fp)
fpr=fp/(fp+tn)
return accuracy,precision,recall,f1,specificity,fpr,cm

start=time.time()
knn_final.fit(X_train_scaled,y_train)
train_time=time.time()-start

start=time.time()
y_pred_knn=knn_final.predict(X_test_scaled)
pred_time=time.time()-start
acc,prec,rec,f1,spec,fpr,cm=compute_metrics(y_test,y_pred_knn)
print("Final KNN Metrics")
print("Accuracy:",acc)
print("Precision:",prec)
print("Recall:",rec)
print("F1 Score:",f1)
print("Specificity:",spec)
print("False Positive Rate:",fpr)
print("Training Time:",train_time)
print("Prediction Time:",pred_time)

```

```

Final KNN Metrics
Accuracy: 0.9207383279044516
Precision: 0.9420731707317073
Recall: 0.8512396694214877
F1 Score: 0.894356005788712
Specificity: 0.9659498207885304
False Positive Rate: 0.034050179211469536
Training Time: 0.0014691352844238281
Prediction Time: 0.03354024887084961

```

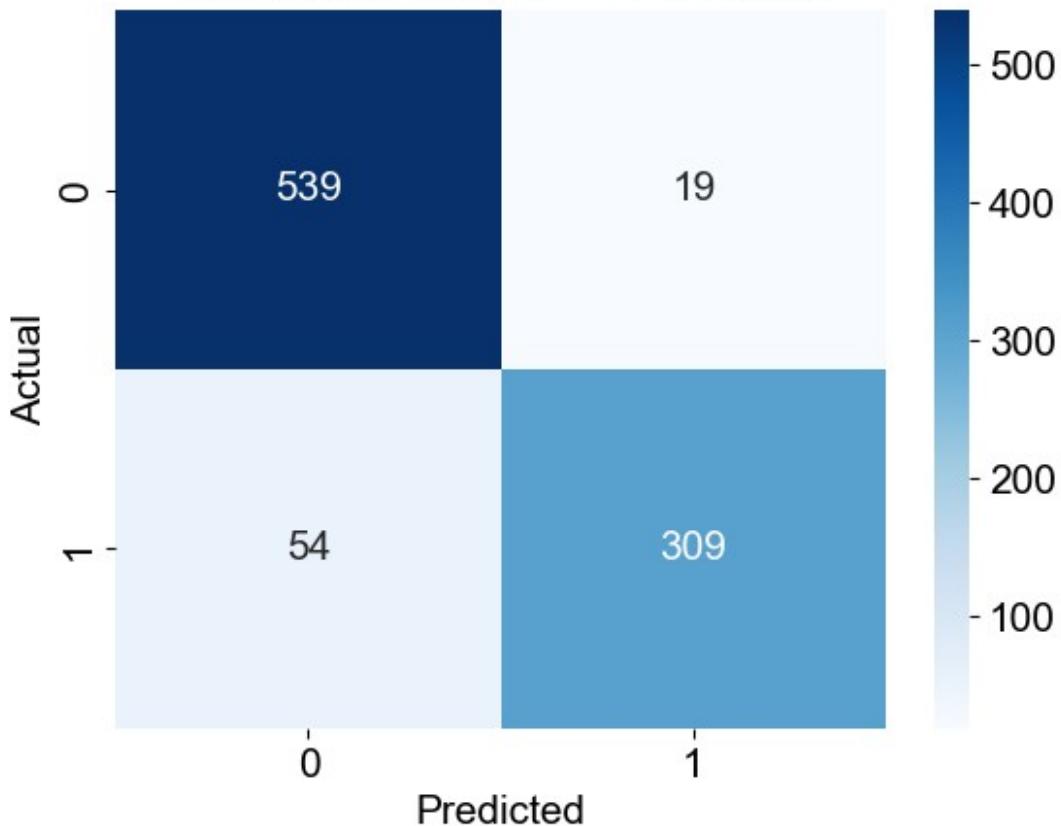
### Confusion Matrix

```

sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.title("Confusion Matrix - Final KNN")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

```

### Confusion Matrix – Final KNN



```
from sklearn.metrics import roc_curve, auc
model_specs = [
    ("Gaussian NB", gnb, X_test),
    ("Bernoulli NB", bnb, X_test),
    ("Multinomial NB", mnb, X_test),
    ("KNN (Best)", knn_final, X_test)
]
plt.figure(figsize=(8, 6))
ax = plt.gca()

colors = ['#1f77b4', '#ff7f0e', '#2ca02c', '#d62728']

for i, (name, model, data) in enumerate(model_specs):
    y_prob = model.predict_proba(data)[:, 1]
    fpr, tpr, _ = roc_curve(y_test, y_prob)
    roc_auc = auc(fpr, tpr)

    ax.plot(fpr, tpr, color=colors[i], lw=2, label=f'{name} (AUC = {roc_auc:.2f})')

ax.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
```

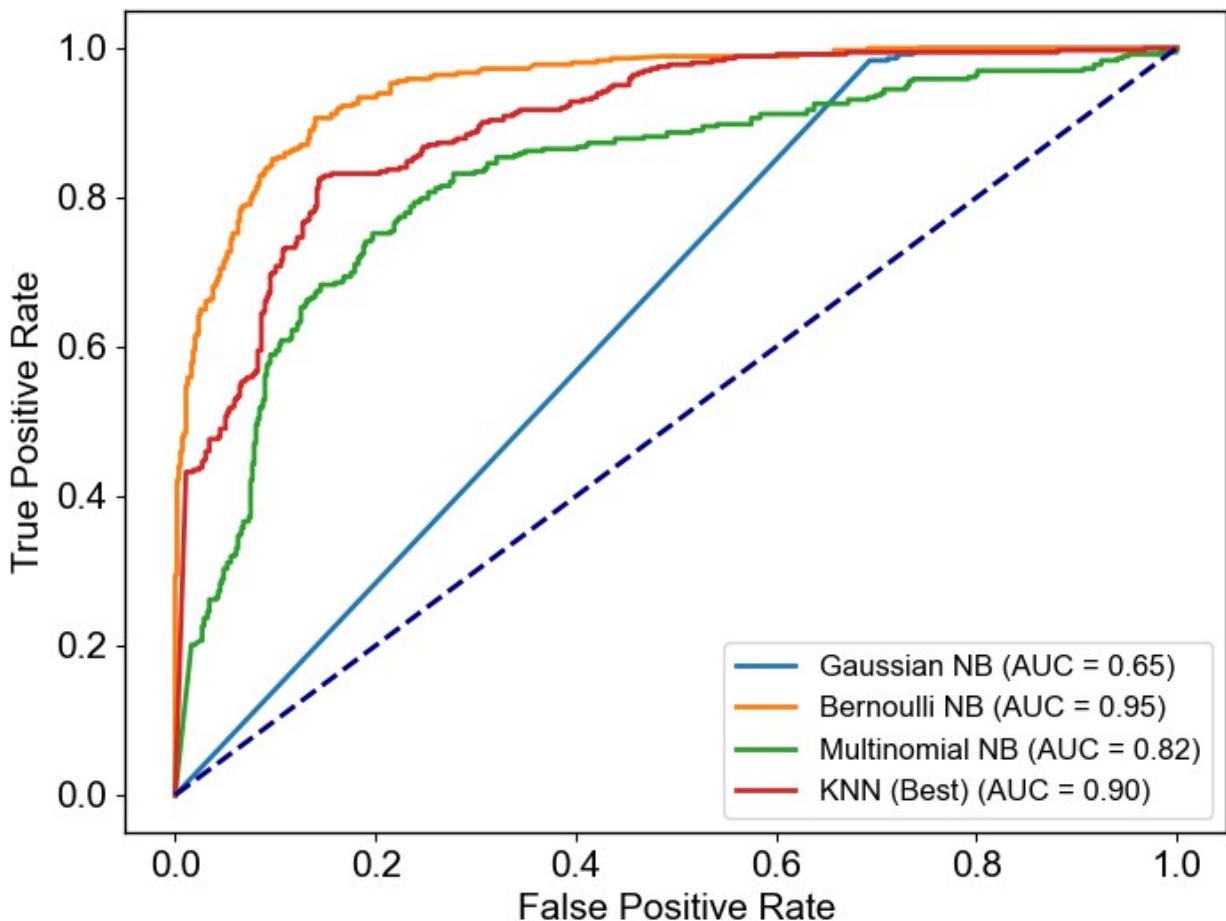
```
ax.set_title("")
ax.set_xlabel('False Positive Rate')
ax.set_ylabel('True Positive Rate')
ax.legend(loc="lower right", prop={'size': 12})

for spine in ax.spines.values():
    spine.set_linewidth(bxwidth)

plt.gcf().text(
    0.5, -0.05,
    "ROC Curves showing the diagnostic ability of each classifier.",
    ha="center", fontsize=15
)

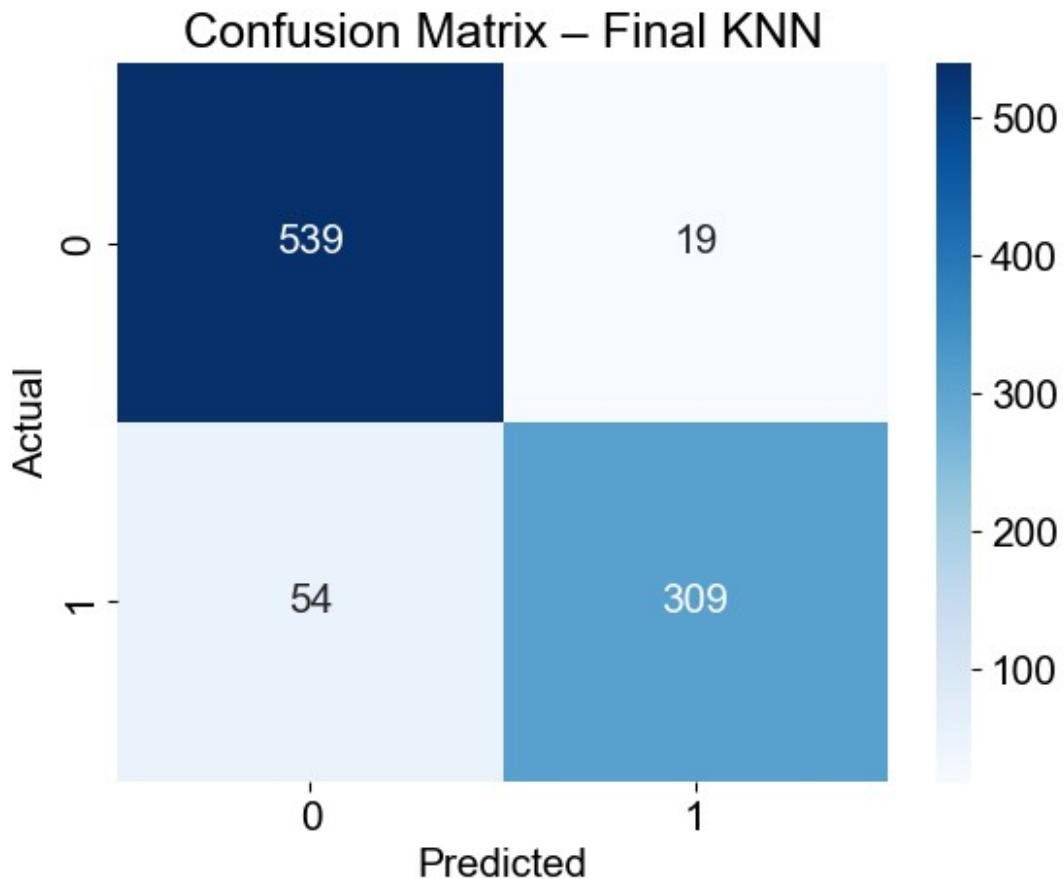
plt.savefig("roc_curves.png", format="png", dpi=600,
bbox_inches="tight")
plt.show()

e:\anaconda\Lib\site-packages\sklearn\utils\validation.py:2732:
UserWarning: X has feature names, but GaussianNB was fitted without
feature names
    warnings.warn(
e:\anaconda\Lib\site-packages\sklearn\utils\validation.py:2732:
UserWarning: X has feature names, but KNeighborsClassifier was fitted
without feature names
    warnings.warn(
```



ROC Curves showing the diagnostic ability of each classifier.

```
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix – Final KNN")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



### KDTree

```

print("---- Step 9: Training Optimized KNN Models (KDTree & BallTree)
----")

best_params=grid.best_params_
optimal_k=best_params['n_neighbors']

start=time.time()
knn_kd=KNeighborsClassifier(
    n_neighbors=optimal_k,
    weights=best_params['weights'],
    metric=best_params['metric'],
    algorithm='kd_tree'
)
knn_kd.fit(X_train_scaled,y_train)
train_time_kd=time.time()-start
start=time.time()
y_pred_kd=knn_kd.predict(X_test_scaled)
pred_time_kd=time.time()-start

acc_kd,prec_kd,rec_kd,f1_kd,sp,fpr,cm=compute_metrics(y_test,y_pred_kd)

```

```

)
print("Final KDtree Metrics")
print("Accuracy:",acc_kd)
print("Precision:",prec_kd)
print("Recall:",rec_kd)
print("F1 Score:",f1_kd)
print("Specificity:",sp)
print("False Positive Rate:",fpr)
print("Training Time:",train_time)
print("Prediction Time:",pred_time)

--- Step 9: Training Optimized KNN Models (KDTree & BallTree) ---
Final KDtree Metrics
Accuracy: 0.9207383279044516
Precision: 0.9420731707317073
Recall: 0.8512396694214877
F1 Score: 0.894356005788712
Specificity: 0.9659498207885304
False Positive Rate: 0.034050179211469536
Training Time: 0.0014691352844238281
Prediction Time: 0.03354024887084961

start=time.time()
knn_bt=KNeighborsClassifier(
    n_neighbors=optimal_k,
    weights=best_params['weights'],
    metric=best_params['metric'],
    algorithm='ball_tree'
)
knn_bt.fit(X_train_scaled,y_train)
train_time_bt=time.time()-start
start=time.time()
y_pred_bt=knn_bt.predict(X_test_scaled)
pred_time_bt=time.time()-start

acc_bt,prec_bt,rec_bt,f1_bt,sp,fpr,cm=compute_metrics(y_test,y_pred_bt)
)
print("Final Balltree Metrics")
print("Accuracy:",acc_bt)
print("Precision:",prec_bt)
print("Recall:",rec_bt)
print("F1 Score:",f1_bt)
print("Specificity:",sp)
print("False Positive Rate:",fpr)
print("Training Time:",train_time)
print("Prediction Time:",pred_time)

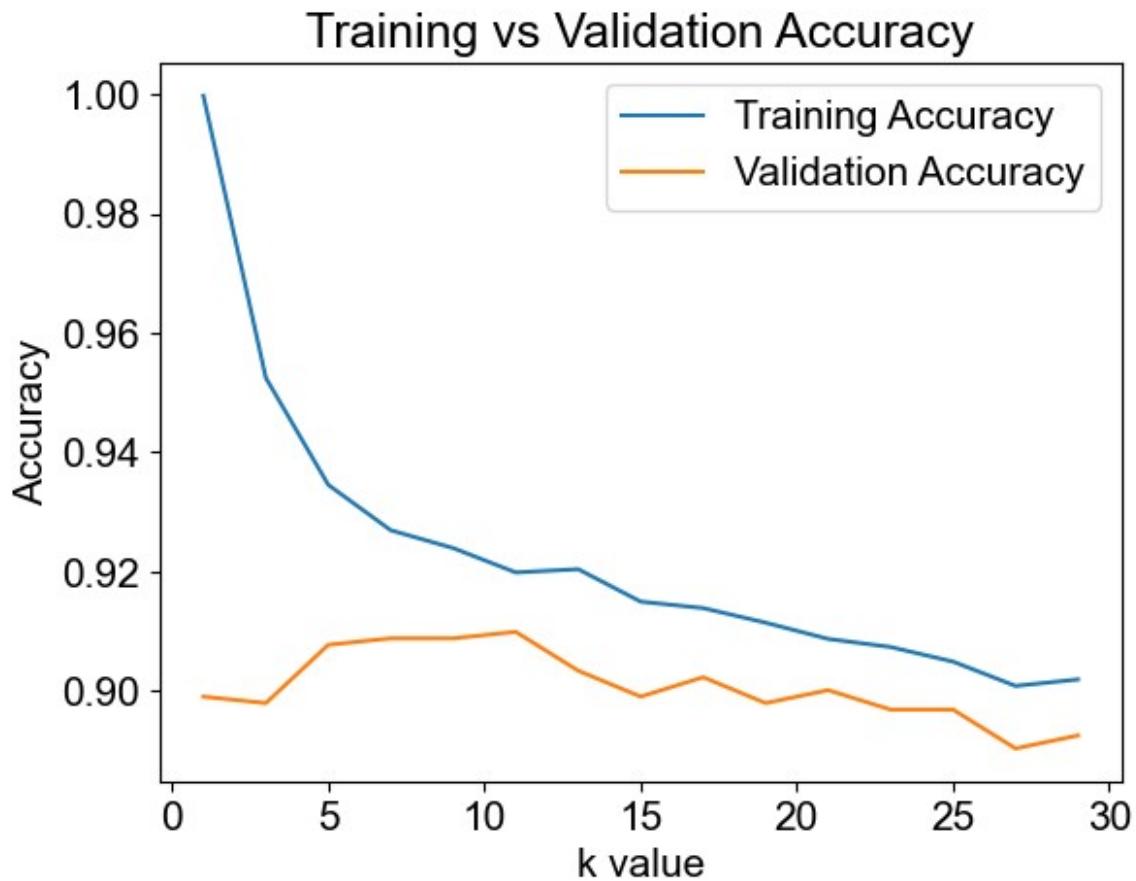
Final Balltree Metrics
Accuracy: 0.9207383279044516
Precision: 0.9420731707317073

```

```
Recall: 0.8512396694214877
F1 Score: 0.894356005788712
Specificity: 0.9659498207885304
False Positive Rate: 0.034050179211469536
Training Time: 0.0014691352844238281
Prediction Time: 0.03354024887084961

train_acc=[]
val_acc=[]
k_values=range(1,31,2)
for k in k_values:
    knn=KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train_scaled,y_train)
    train_acc.append(knn.score(X_train_scaled,y_train))
    val_acc.append(knn.score(X_test_scaled,y_test))
plt.plot(k_values,train_acc,label='Training Accuracy')
plt.plot(k_values,val_acc,label='Validation Accuracy')
plt.xlabel('k value')
plt.ylabel('Accuracy')
plt.title('Training vs Validation Accuracy')
plt.legend()
plt.savefig("train_valid_validation.png",dpi = 600 , bbox_inches =
"tight")

plt.show()
```



### Cross Validation Scores

```
gnb_scores=cross_val_score(  
    GaussianNB(),  
    X_train,  
    y_train,  
    cv=skf,  
    scoring='accuracy'  
)  
mnb_scores=cross_val_score(  
    MultinomialNB(),  
    X_train,  
    y_train,  
    cv=skf,  
    scoring='accuracy'  
)  
bnb_scores=cross_val_score(  
    BernoulliNB(),  
    X_train,  
    y_train,  
    cv=skf,  
    scoring='accuracy'
```

```

)
knn_kd_scores=cross_val_score(
    knn_kd,
    X_train_scaled,
    y_train,
    cv=skf,
    scoring='accuracy'
)
knn_bt_scores=cross_val_score(
    knn_bt,
    X_train_scaled,
    y_train,
    cv=skf,
    scoring='accuracy'
)

print("One Way test")
from scipy.stats import f_oneway
F_stat,p_value=f_oneway(
    gnb_scores,
    mnb_scores,
    bnb_scores,
    knn_kd_scores,
    knn_bt_scores
)

print("F-statistic:",F_stat)
print("p-value:",p_value)
print("Gaussian NB Mean Accuracy:",gnb_scores.mean())
print("Multinomial NB Mean Accuracy:",mnb_scores.mean())
print("Bernoulli NB Mean Accuracy:",bnb_scores.mean())
print("KNN KDTree Mean Accuracy:",knn_kd_scores.mean())
print("KNN BallTree Mean Accuracy:",knn_bt_scores.mean()
))
best_model=max(
    [
        ("Gaussian NB",gnb_scores.mean()),
        ("Multinomial NB",mnb_scores.mean()),
        ("Bernoulli NB",bnb_scores.mean()),
        ("KNN KDTree",knn_kd_scores.mean()),
        ("KNN BallTree",knn_bt_scores.mean())
    ],
    key=lambda x:x[1]
)
print("Best Model:",best_model[0])

One Way test
F-statistic: 83.72948758252099
p-value: 3.3695720596139815e-12

```

```
Gaussian NB Mean Accuracy: 0.8206521739130436
Multinomial NB Mean Accuracy: 0.7869565217391304
Bernoulli NB Mean Accuracy: 0.8872282608695652
KNN KDTree Mean Accuracy: 0.9252717391304348
KNN BallTree Mean Accuracy: 0.9252717391304348
Best Model: KNN KDTree
```