

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import learning_curve
df = pd.read_csv("E:/ml/train.csv")

df.head()

   Customer ID          Name Gender  Age Income (USD) Income Stability \
0      C-36995  Frederica Shealy      F   56     1933.05      Low
1      C-33999    America Calderone      M   32     4952.91      Low
2      C-3770       Rosetta Verne      F   65      988.19     High
3      C-26480        Zoe Chitty      F   65        NaN     High
4      C-23459       Afton Venema      F   31     2614.77      Low

   Profession Type of Employment Location  Loan Amount Request \
0    Working           Sales staff  Semi-Urban 72809.58
1    Working            NaN  Semi-Urban 46837.47
2 Pensioner            NaN  Semi-Urban 45593.04
3 Pensioner            NaN      Rural 80057.92
4    Working  High skill tech staff  Semi-Urban 113858.89

   ...  Credit Score No. of Defaults Has Active Credit Card  Property \
ID \
0 ...      809.44                  0                NaN
746 ...
1 ...      780.40                  0        Unpossessed
608 ...
2 ...      833.15                  0        Unpossessed
546 ...
3 ...      832.70                  1        Unpossessed
890 ...
4 ...      745.55                  1         Active
715 ...

```

```

Property Age  Property Type  Property Location  Co-Applicant \
0      1933.05          4            Rural           1
1      4952.91          2            Rural           1
2      988.19          2            Urban           0
3        NaN          2  Semi-Urban           1
4     2614.77          4  Semi-Urban           1

```

```

Property Price  Loan Sanction Amount (USD)
0      119933.46          54607.18
1      54791.00          37469.98
2      72440.58          36474.43
3      121441.51          56040.54
4      208567.91          74008.28

```

[5 rows x 24 columns]

```
df.drop(['Customer ID', 'Name', 'Property ID'], axis =1 ,inplace=True)
```

```
print(df.shape)
```

```
df = df.dropna(subset=['Loan Sanction Amount (USD)'])
```

```
print (df.shape)
```

```
(30000, 21)
(29660, 21)
```

```
missing_counts = df.isnull().sum()
```

```
print(missing_counts)
```

Gender	52
Age	0
Income (USD)	4493
Income Stability	1658
Profession	0
Type of Employment	7188
Location	0
Loan Amount Request (USD)	0
Current Loan Expenses (USD)	167
Expense Type 1	0
Expense Type 2	0
Dependents	2446
Credit Score	1670
No. of Defaults	0
Has Active Credit Card	1546
Property Age	4760
Property Type	0
Property Location	347
Co-Applicant	0

```

Property Price          0
Loan Sanction Amount (USD)    0
dtype: int64

num_cols = df.select_dtypes(include=['int64', 'float64']).columns
cat_cols = df.select_dtypes(include=['object']).columns

for col in num_cols:
    df[col] = df[col].fillna(df[col].median())

encoders = {}

for col in cat_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col].astype(str))
    encoders[col] = le
print("Missing values after handling [VERIFICATION]")
print(df.isnull().sum().sum())

Missing values after handling [VERIFICATION]
0

le = LabelEncoder()

for col in cat_cols:
    df[col] = le.fit_transform(df[col])

#Dropping loan amount 0
df = df[df['Loan Sanction Amount (USD)'] > 0]

print(df.describe())

      Gender        Age   Income (USD)  Income Stability \
count  21457.000000  21457.000000  2.145700e+04  21457.000000
mean    0.504591    40.389430  2.608860e+03  0.966957
std     0.503428    16.212181  1.222496e+04  0.407881
min     0.000000    18.000000  3.787600e+02  0.000000
25%    0.000000    25.000000  1.769020e+03  1.000000
50%    1.000000    40.000000  2.223300e+03  1.000000
75%    1.000000    56.000000  2.869780e+03  1.000000
max    2.000000    65.000000  1.777460e+06  2.000000

      Profession  Type of Employment  Location \
count  21457.000000  21457.000000  21457.000000
mean    4.711516    10.483991   0.940159
std     2.626624     5.922649   0.525947
min     0.000000     0.000000   0.000000
25%    1.000000     6.000000   1.000000
50%    7.000000    10.000000   1.000000
75%    7.000000    18.000000   1.000000

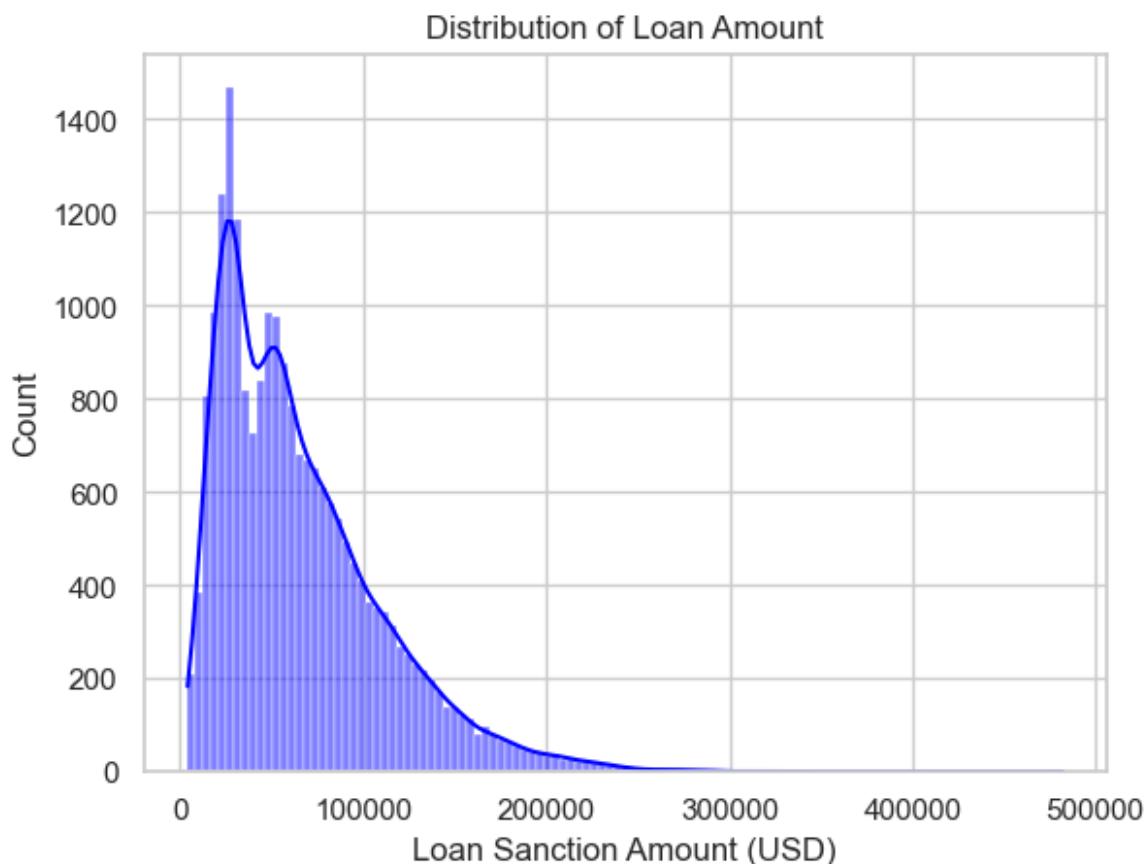
```

max	7.000000	18.000000	2.000000
Loan Amount Request (USD) Current Loan Expenses (USD) Expense			
Type	1 \		
count	21457.000000		21457.000000
mean	91495.926906		405.901925
std	60796.889284		249.574325
min	6108.050000		-999.000000
0.000000			
25%	42053.750000		251.630000
0.000000			
50%	77334.440000		378.320000
0.000000			
75%	123930.540000		527.810000
1.000000			
max	602384.150000		3840.880000
1.000000			
Dependents Credit Score No. of Defaults \			
count	21457.000000	21457.000000	21457.000000
mean	2.227758	756.849588	0.193690
std	0.893101	60.774241	0.395198
min	1.000000	620.080000	0.000000
25%	2.000000	708.550000	0.000000
50%	2.000000	749.060000	0.000000
75%	3.000000	805.920000	0.000000
max	10.000000	896.260000	1.000000
Has Active Credit Card Property Age Property Type Property			
Location \			
count	21457.000000	2.145700e+04	21457.000000
21457.000000			
mean	1.089481	2.604765e+03	2.458825
1.000233			
std	0.912621	1.222387e+04	1.119531
0.835372			
min	0.000000	3.787600e+02	1.000000
0.000000			
25%	0.000000	1.787940e+03	1.000000
0.000000			
50%	1.000000	2.223965e+03	2.000000
1.000000			
75%	2.000000	2.849640e+03	3.000000
2.000000			
max	3.000000	1.777460e+06	4.000000
3.000000			

	Co-Applicant	Property Price	Loan Sanction Amount (USD)
count	21457.000000	2.145700e+04	21457.000000
mean	-5.247891	1.355782e+05	65881.397768
std	78.484035	9.573963e+04	44858.650085
min	-999.000000	-9.990000e+02	4023.180000
25%	1.000000	6.245823e+04	30310.860000
50%	1.000000	1.133293e+05	55191.660000
75%	1.000000	1.852552e+05	89126.820000
max	1.000000	1.028083e+06	481907.320000

[8 rows x 21 columns]

```
sns.histplot(df['Loan Sanction Amount (USD)'], kde=True, color='blue')
plt.title('Distribution of Loan Amount')
plt.show()
```



```
fig, axes = plt.subplots(1, 3, figsize=(18, 6))

sns.countplot(x='Income Stability', data=df, ax=axes[0],
palette='viridis')
axes[0].set_title('Income Stability Count')

sns.countplot(x='Profession', data=df, ax=axes[1], palette='viridis')
```

```
axes[1].set_title('Profession Count')

sns.countplot(x='Location', data=df, ax=axes[2], palette='magma')
axes[2].set_title('Location Count')

plt.tight_layout()
plt.show()

print("--- Income Stability Mapping ---")
for index, label in enumerate(encoders['Income Stability'].classes_):
    print(f"Number {index} means: {label}")

print("\n--- Profession Mapping ---")
for index, label in enumerate(encoders['Profession'].classes_):
    print(f"Number {index} means: {label}")

print("\n--- Location Mapping ---")
for index, label in enumerate(encoders['Location'].classes_):
    print(f"Number {index} means: {label}")

C:\Users\KESHA\AppData\Local\Temp\ipykernel_27448\1852018604.py:3:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.

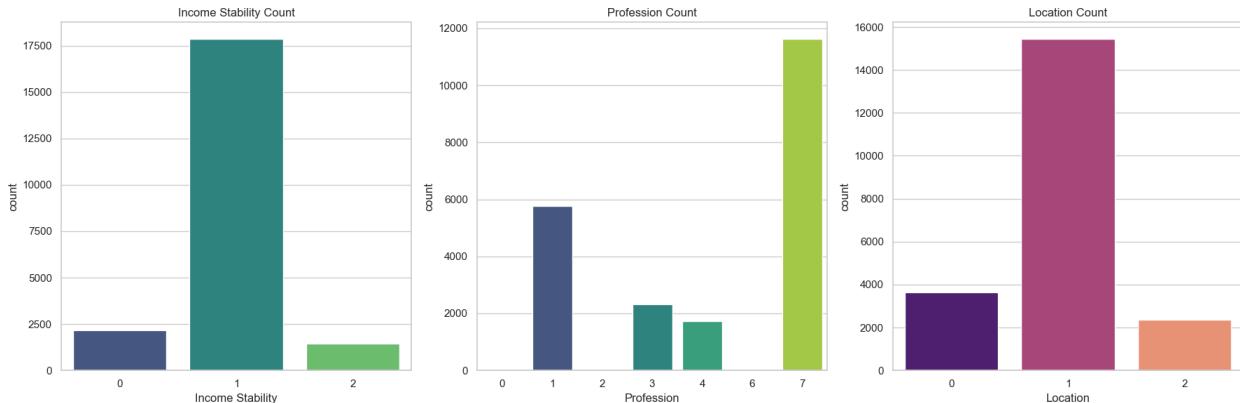
    sns.countplot(x='Income Stability', data=df, ax=axes[0],
palette='viridis')
C:\Users\KESHA\AppData\Local\Temp\ipykernel_27448\1852018604.py:6:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.

    sns.countplot(x='Profession', data=df, ax=axes[1],
palette='viridis')
C:\Users\KESHA\AppData\Local\Temp\ipykernel_27448\1852018604.py:10:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.

    sns.countplot(x='Location', data=df, ax=axes[2], palette='magma')
```



```
--- Income Stability Mapping ---
```

Number 0 means: High

Number 1 means: Low

Number 2 means: nan

```
--- Profession Mapping ---
```

Number 0 means: Businessman

Number 1 means: Commercial associate

Number 2 means: Maternity leave

Number 3 means: Pensioner

Number 4 means: State servant

Number 5 means: Student

Number 6 means: Unemployed

Number 7 means: Working

```
--- Location Mapping ---
```

Number 0 means: Rural

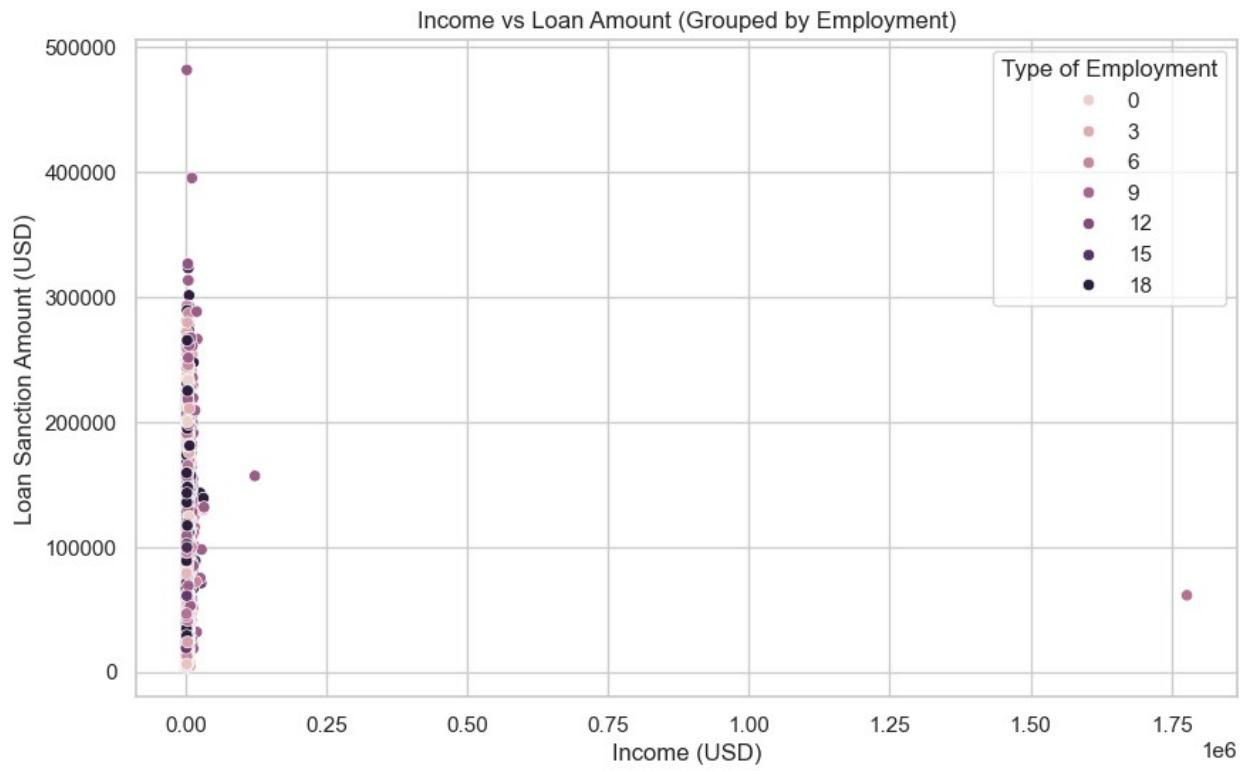
Number 1 means: Semi-Urban

Number 2 means: Urban

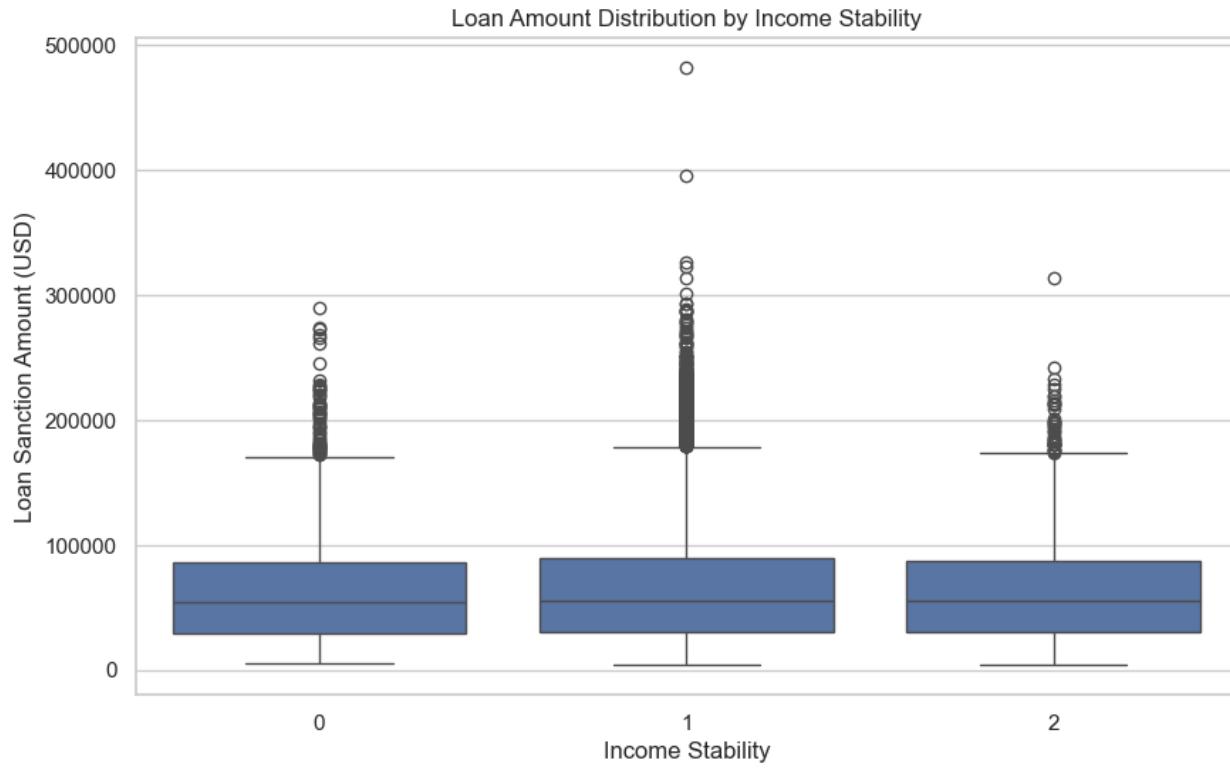
```
plt.figure(figsize=(10, 6))
sns.scatterplot(x='Income (USD)', y='Loan Sanction Amount (USD)',
hue='Type of Employment', data=df)# use legend='full' if want to see
the full thing
plt.title('Income vs Loan Amount (Grouped by Employment)')
plt.show()
```

```
for index, label in enumerate(encoders['Type of
Employment'].classes_):
    print(f"Number {index} means: {label}")
```

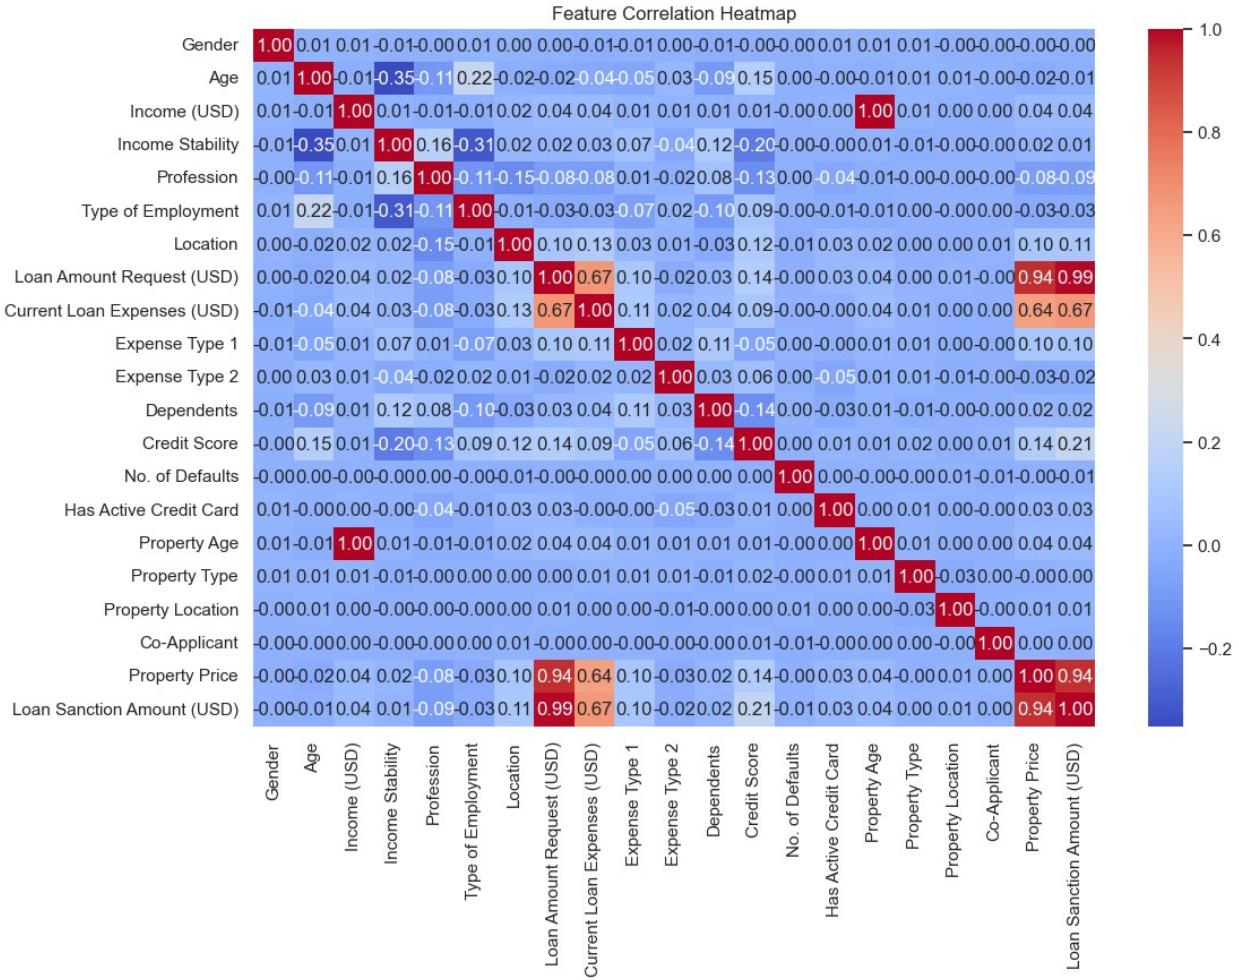
```
plt.figure(figsize=(10, 6))
sns.boxplot(x='Income Stability', y='Loan Sanction Amount (USD)',
data=df)
plt.title('Loan Amount Distribution by Income Stability')
plt.show()
```



Number 0 means: Accountants  
 Number 1 means: Cleaning staff  
 Number 2 means: Cooking staff  
 Number 3 means: Core staff  
 Number 4 means: Drivers  
 Number 5 means: HR staff  
 Number 6 means: High skill tech staff  
 Number 7 means: IT staff  
 Number 8 means: Laborers  
 Number 9 means: Low-skill Laborers  
 Number 10 means: Managers  
 Number 11 means: Medicine staff  
 Number 12 means: Private service staff  
 Number 13 means: Realty agents  
 Number 14 means: Sales staff  
 Number 15 means: Secretaries  
 Number 16 means: Security staff  
 Number 17 means: Waiters/barmen staff  
 Number 18 means: nan



```
plt.figure(figsize=(12, 8))
correlation_matrix = df.select_dtypes(include=[np.number]).corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
            fmt='.2f')
plt.title('Feature Correlation Heatmap')
plt.show()
```



```

from sklearn.model_selection import train_test_split

X = df.drop('Loan Sanction Amount (USD)', axis=1)
y = df['Loan Sanction Amount (USD)']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_final = pd.DataFrame(X_scaled, columns=X.columns)
X = X_final

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

print(f"Training shapes: {X_train.shape}, {y_train.shape}")
print(f"Testing shapes: {X_test.shape}, {y_test.shape}")

```

```

Training shapes: (17165, 20), (17165,)
Testing shapes: (4292, 20), (4292,)

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score

model = LinearRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print(f"Mean Absolute Error: {mae:.2f}")
print(f"RMSE: {rmse:.2f}")
print(f"R2 Score: {r2:.2f}")

Mean Absolute Error: 3736.34
RMSE: 5576.15
R2 Score: 0.99

from sklearn.linear_model import Ridge, Lasso, ElasticNet
from sklearn.metrics import r2_score, mean_absolute_error

ridge = Ridge(alpha=0.1)
ridge.fit(X_train, y_train)
y_pred_ridge = ridge.predict(X_test)

lasso = Lasso(alpha=0.1)
lasso.fit(X_train, y_train)
y_pred_lasso = lasso.predict(X_test)

enet = ElasticNet(alpha=0.1, l1_ratio=0.5) # l1_ratio=0.5 means 50%
# Lasso, 50% Ridge
enet.fit(X_train, y_train)
y_pred_enet = enet.predict(X_test)

results = {
    "Model": ["Ridge", "Lasso", "Elastic Net"],
    "MAE": [
        mean_absolute_error(y_test, y_pred_ridge),
        mean_absolute_error(y_test, y_pred_lasso),
        mean_absolute_error(y_test, y_pred_enet)
    ],
    "R2 Score": [
        r2_score(y_test, y_pred_ridge),
        r2_score(y_test, y_pred_lasso),

```

```

        r2_score(y_test, y_pred_enet)
    ]
}

print(pd.DataFrame(results))

      Model      MAE  R2 Score
0    Ridge  3736.315081  0.985198
1   Lasso  3735.974004  0.985205
2 Elastic Net  4460.673938  0.978412

from sklearn.model_selection import GridSearchCV

param_grid = {'alpha': [0.001, 0.01, 0.1, 1, 10, 100]}

ridge_cv = GridSearchCV(Ridge(), param_grid, cv=5, scoring='r2')
ridge_cv.fit(X_train, y_train)

lasso_cv = GridSearchCV(Lasso(), param_grid, cv=5, scoring='r2')
lasso_cv.fit(X_train, y_train)

print(f"Best Ridge Alpha: {ridge_cv.best_params_}")
print(f"Best Lasso Alpha: {lasso_cv.best_params_}")

enet_params = {
    'alpha': [0.1, 1, 10],
    'l1_ratio': [0.2, 0.5, 0.8]
}

enet_cv = GridSearchCV(ElasticNet(), enet_params, cv=5, scoring='r2')
enet_cv.fit(X_train, y_train)

print(f"Best Elastic Net Params: {enet_cv.best_params_}")

e:\anaconda\Lib\site-packages\sklearn\linear_model\
_coordinate_descent.py:695: ConvergenceWarning: Objective did not
converge. You might want to increase the number of iterations, check
the scale of the features or consider increasing regularisation.
Duality gap: 5.587e+10, tolerance: 2.756e+09
    model = cd_fast.enet_coordinate_descent(
e:\anaconda\Lib\site-packages\sklearn\linear_model\
_coordinate_descent.py:695: ConvergenceWarning: Objective did not
converge. You might want to increase the number of iterations, check
the scale of the features or consider increasing regularisation.
Duality gap: 7.446e+10, tolerance: 2.713e+09
    model = cd_fast.enet_coordinate_descent(
e:\anaconda\Lib\site-packages\sklearn\linear_model\
_coordinate_descent.py:695: ConvergenceWarning: Objective did not
converge. You might want to increase the number of iterations, check
the scale of the features or consider increasing regularisation.
Duality gap: 1.434e+11, tolerance: 2.742e+09

```

```

model = cd_fast.enet_coordinate_descent(
e:\anaconda\Lib\site-packages\sklearn\linear_model\
_coordinate_descent.py:695: ConvergenceWarning: Objective did not
converge. You might want to increase the number of iterations, check
the scale of the features or consider increasing regularisation.
Duality gap: 2.560e+10, tolerance: 2.743e+09
    model = cd_fast.enet_coordinate_descent(
e:\anaconda\Lib\site-packages\sklearn\linear_model\
_coordinate_descent.py:695: ConvergenceWarning: Objective did not
converge. You might want to increase the number of iterations, check
the scale of the features or consider increasing regularisation.
Duality gap: 4.647e+10, tolerance: 2.707e+09
    model = cd_fast.enet_coordinate_descent(
e:\anaconda\Lib\site-packages\sklearn\linear_model\
_coordinate_descent.py:695: ConvergenceWarning: Objective did not
converge. You might want to increase the number of iterations, check
the scale of the features or consider increasing regularisation.
Duality gap: 3.918e+09, tolerance: 2.713e+09
    model = cd_fast.enet_coordinate_descent(
e:\anaconda\Lib\site-packages\sklearn\linear_model\
_coordinate_descent.py:695: ConvergenceWarning: Objective did not
converge. You might want to increase the number of iterations, check
the scale of the features or consider increasing regularisation.
Duality gap: 2.897e+10, tolerance: 2.742e+09
    model = cd_fast.enet_coordinate_descent()

Best Ridge Alpha: {'alpha': 1}
Best Lasso Alpha: {'alpha': 10}
Best Elastic Net Params: {'alpha': 0.1, 'l1_ratio': 0.8}

best_ridge = ridge_cv.best_estimator_
best_lasso = lasso_cv.best_estimator_
best_enet = enet_cv.best_estimator_

models = [best_ridge, best_lasso, best_enet]
model_names = ["Tuned Ridge", "Tuned Lasso", "Tuned Elastic Net"]

for name, model in zip(model_names, models):
    score = model.score(X_test, y_test)
    print(f"{name} R2 Score: {score:.4f}")

Tuned Ridge R2 Score: 0.9852
Tuned Lasso R2 Score: 0.9856
Tuned Elastic Net R2 Score: 0.9834

from sklearn.model_selection import cross_validate
from sklearn.metrics import make_scorer, mean_absolute_error,
mean_squared_error, r2_score
import numpy as np

```

```

models = {
    "Linear Regression": LinearRegression(),
    "Ridge Regression": ridge_cv.best_estimator_,
    "Lasso Regression": lasso_cv.best_estimator_,
    "Elastic Net Regression": enet_cv.best_estimator_
}

scoring = {
    "mae": make_scorer(mean_absolute_error, greater_is_better=False),
    "mse": make_scorer(mean_squared_error, greater_is_better=False),
    "r2": make_scorer(r2_score)
}

cv_results = []

for name, model in models.items():
    scores = cross_validate(model, X_train, y_train, cv=5,
                           scoring=scoring)
    cv_results.append([
        name,
        -scores["test_mae"].mean(),
        -scores["test_mse"].mean(),
        np.sqrt(-scores["test_mse"].mean()),
        scores["test_r2"].mean()
    ])

table2 = pd.DataFrame(
    cv_results,
    columns=["Model", "MAE", "MSE", "RMSE", "R2"]
)

```

table2

	Model	MAE	MSE	RMSE
R <sup>2</sup>				
0	Linear Regression	3609.000583	2.764272e+07	5257.634035
0.986109				
1	Ridge Regression	3608.588185	2.764016e+07	5257.391122
0.986110				
2	Lasso Regression	3605.700848	2.762616e+07	5256.059731
0.986117				
3	Elastic Net Regression	3825.689632	3.179339e+07	5638.563133
0.984025				

```

models = {
    "Linear Regression": model,
    "Tuned Ridge": best_ridge,
    "Tuned Lasso": best_lasso,
    "Tuned Elastic Net": best_enet
}

```

```

eval_metrics = []

for name, model_obj in models.items():
    preds = model_obj.predict(X_test)

    mae = mean_absolute_error(y_test, preds)
    mse = mean_squared_error(y_test, preds)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_test, preds)

    eval_metrics.append({
        "Model": name,
        "MAE": round(mae, 2),
        "MSE": round(mse, 2),
        "RMSE": round(rmse, 2),
        "R2 Score": round(r2, 4)
    })

results_df = pd.DataFrame(eval_metrics)
print(results_df.sort_values(by="R2 Score", ascending=False))

```

	Model	MAE	MSE	RMSE	R2 Score
2	Tuned Lasso	3719.83	30323919.93	5506.72	0.9856
1	Tuned Ridge	3736.13	31082233.67	5575.14	0.9852
0	Linear Regression	3960.60	34931250.26	5910.27	0.9834
3	Tuned Elastic Net	3960.60	34931250.26	5910.27	0.9834

```

coef_table = pd.DataFrame({
    "Feature": X_train.columns,
    "Linear": models["Linear Regression"].coef_,
    "Ridge": models["Tuned Ridge"].coef_,
    "Lasso": models["Tuned Lasso"].coef_,
    "Elastic Net": models["Tuned Elastic Net"].coef_
})

```

coef\_table

	Feature	Linear	Ridge
Lasso \ 0	Gender	17.580328	-13.749662
3.332303			-
1	Age	2.577235	-8.053297
0.000000			-
2	Income (USD)	14.711576	467.419995
0.000000			-
3	Income Stability	113.774970	96.178853
83.734410			-
4	Profession	-62.116131	-49.089305
34.995102			-

5	Type of Employment	-95.298883	-77.458892	-
67.525299				
6	Location	64.721397	72.108777	
59.081410				
7	Loan Amount Request (USD)	37114.264323	43555.327303	
43567.138220				
8	Current Loan Expenses (USD)	724.552055	-105.200195	-
87.614799				
9	Expense Type 1	1.649791	-40.258424	-
35.701416				
10	Expense Type 2	11.295738	31.770559	
19.856662				
11	Dependents	-58.948498	-62.172947	-
51.819072				
12	Credit Score	3030.738136	2989.115500	
2980.104096				
13	No. of Defaults	-54.441014	-42.373181	-
32.993159				
14	Has Active Credit Card	81.524001	46.118142	
36.565360				
15	Property Age	0.559413	-858.754419	-
0.000000				
16	Property Type	96.813176	92.930411	
82.030963				
17	Property Location	27.391578	1.153002	
0.000000				
18	Co-Applicant	-14.548111	-6.096593	-
0.000000				
19	Property Price	5968.459789	545.564251	
502.370322				

#### Elastic Net

0	17.580328
1	2.577235
2	14.711576
3	113.774970
4	-62.116131
5	-95.298883
6	64.721397
7	37114.264323
8	724.552055
9	1.649791
10	11.295738
11	-58.948498
12	3030.738136
13	-54.441014
14	81.524001
15	0.559413
16	96.813176

```

17      27.391578
18     -14.548111
19     5968.459789

y_pred = best_ridge.predict(X_test)
residuals = y_test - y_pred

sns.set_theme(style="whitegrid")
fig, axes = plt.subplots(3, 2, figsize=(16, 18))

sns.histplot(y_test, kde=True, ax=axes[0, 0], color='teal')
axes[0, 0].set_title('Target Distribution: Loan Amount')

sns.scatterplot(x=X_test['Income (USD)'], y=y_test, ax=axes[0, 1],
alpha=0.6)
axes[0, 1].set_title('Applicant Income vs. Loan Amount')

sns.scatterplot(x=y_test, y=y_pred, ax=axes[1, 0], color='blue',
alpha=0.5)
axes[1, 0].plot([y_test.min(), y_test.max()], [y_test.min(),
y_test.max()], 'r--', lw=2)
axes[1, 0].set_title('Predicted vs. Actual Values')
axes[1, 0].set_xlabel('Actual')
axes[1, 0].set_ylabel('Predicted')

sns.scatterplot(x=y_pred, y=residuals, ax=axes[1, 1], color='purple',
alpha=0.5)
axes[1, 1].axhline(0, color='red', linestyle='--')
axes[1, 1].set_title('Residual Plot (Errors vs. Predictions)')
axes[1, 1].set_xlabel('Predicted Values')
axes[1, 1].set_ylabel('Residuals')

train_sizes, train_scores, val_scores = learning_curve(best_ridge, X,
y, cv=5, scoring='neg_mean_absolute_error')
train_mae = -train_scores.mean(axis=1)
val_mae = -val_scores.mean(axis=1)

axes[2, 0].plot(train_sizes, train_mae, label='Training Error')
axes[2, 0].plot(train_sizes, val_mae, label='Validation Error')
axes[2, 0].set_title('Learning Curve: Training vs. Validation Error')
axes[2, 0].legend()

coefs = pd.Series(best_ridge.coef_, index=X.columns).sort_values()
coefs.plot(kind='barh', ax=axes[2, 1], color='skyblue')
axes[2, 1].set_title('Feature Importance (Model Coefficients)')
plt.savefig('regression_analysis_results.png', dpi=600, format="png",
bbox_inches='tight')
plt.tight_layout()
plt.show()

```

